Title

MARKUS ARONSSON

Department of Funktionell Programmering
CHALMERS UNIVERSITY OF TECHNOLOGY
Gothenburg, Sweden TBD

Title

Markus Aronsson

to my family

Abstract

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

Keywords: Product configuration, constraint satisfaction, Boolean satisfiability, knowledge compilation, supervisory control theory.

Acknowledgments

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Markus Aronsson
Göteborg, December 2012

List of Publications

This thesis is based on the following appended papers:

Paper ??. Markus Aronsson, Emil Axelsson, Mary Sheeran. *Stream Processing for Embedded Domain Specific Languages.* Where?

Paper ??. Markus Aronsson, Mary Sheeran. *Hardware Software Co-design in Haskell.* Where?

Contents

Part I

Introductory chapters

Chapter 1

Introduction

Over the last few years, the amount of traffic going back and forth between devices in the global communications infrastructure has been increasing at a rapid pace. Considering how common connected gadgets have become this might not come as big surprise, but as the steady growth of mobile users and internet of things devices continues (Obile 2016), the amount of mobile traffic is projected to become even greater. In fact, global internet traffic is estimated to grow at an average rate of twenty two percent annually, reaching approximately two hundred and fifty million terabytes per month by the end of 2021 (Networking Index 2016). For the communication infrastructure, the consequence of such a rapid growth rate has been a sharp increase in the demand for computational power on systems that already run under tight latency constraints and with limited memory (Persson 2014), which means computations have to be efficient.

High demands for efficiency under tight latency and resource constraints have greatly influenced the development of embedded systems used for communications infrastructure. Today, digital signal processing software for embedded systems is typically written in low level dialect of C. This choice of language is primarily driven by the desire to access the full potential of a processor or memory system, where the use of high level abstractions could have introduced inefficiencies. Low level languages do however force developers to focus on low level implementation details rather than any higher level, mathematical specification of the algorithm they are developing. This in turn has the unfortunate consequence of discouraging developers from focusing on other, less immediate aspects of software development such as portability and modularity. Its difficult to get back these aspects later in the development cycle, as by then any design decisions will be tightly coupled with the implementation.

As important as it is to increase the computational power of embedded systems, limiting their power consumption presents an equally important issue in their architectural design (Mudge 2001). The trend of trading power for performance cannot continue indefinitely: house hold processors today typically have a power density of 70W and upwards, which is seven times that of a typical hot plate. Containing the growth in power requires architectural improvements, with specialized computing for specialized tasks. Heterogeneous computing represents an interesting development

towards the goal of energy efficient computing, and refers to systems that use more than one kind of processor or cores. These systems gain performance or energy efficiency not just by adding more of the same type of processors, but by adding different kinds of co-processors, usually incorporating specialized processing capabilities to handle particular tasks.

Heterogeneous computing systems present new software challenges not found in the development for typical homogeneous systems: sequential programming of low level signal processing is hard, parallel programming is harder, and parallel programming of heterogeneous machines is a considerable challenge. Though performance and power efficiency constraints further complicates the development of such systems.

This thesis consists of two parts. Part I is a general introduction to the field and puts the appended papers into context. Part II contains the appended papers.

1.1   Background

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

1.2   Functional programming

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

1.3   Domain Specific Languages

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum

ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

1.4 Domain Specific Embedded Languages

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetuer adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

Chapter 2

Co-Design

Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper, leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellentesque, augue quis sagittis posuere, turpis lacus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla facilisi. Sed a turpis eu lacus commodo facilisis. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetuer.

2.1   Section about Expression

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

```
1  square :: SExp Int32 → SExp Int32
2  square a = a * a
```

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat.

Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

```
1  square :: HExp Int32 → HExp Int32
2  square a = a * a
```

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

```
1  square :: (Multiplicative exp, Type' exp a, Num a) ⇒ exp a → exp a
2  square a = a * a
```

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

```
1  type Point a = (a, a)
2
3  pair :: (Expr exp, Type' exp a, Num a) ⇒ Point (exp a) → Point (exp a) → exp a
4  pair (a, b) (u, v) = (a + b) * (u + v)
```

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus.

Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

```
1  dotProd :: (Expr exp, Type' exp a, Num a) ⇒ Pull exp a → Pull exp a → exp a
2  dotProd xs ys = forLoop n 0 $ λi s → s + xs!i * ys!i
3    where
4      n = min (length xs) (length ys)
```

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

```
1  forLoop :: Syntax exp st ⇒ exp Length → st → (exp Index → st → st) → st
```

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

```
1  zipWith :: Expr exp ⇒ (a → b → c) → Pull exp a → Pull exp b → Pull exp c
2  zipWith f xs ys = fmap (uncurry f) (zip xs ys)
3
4  sum :: (Expr exp, Type' exp a, Num a) ⇒ Pull exp a → exp a
5  sum = fold (+) 0
```

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

```
1  scProd :: (Expr exp, Type' exp a, Num a) ⇒ Pull exp a → Pull exp a → exp a
2  scProd a b = sum (zipWith (*) a b)
```

2.2   Section about Programs

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum
ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu
libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu
neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames
ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus
vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat.
Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel
leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar
at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus.
Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit
amet orci dignissim rutrum.

```
1  hello :: Software ()
2  hello = printf "Hello world!\n"
```

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor
lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus.
Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae
ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut
massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et
magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla
ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

```
1  hello :: Software ()
2  hello = printf "Hello world!\n"
```

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at,
tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy
pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac
quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas
lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi
blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae
enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod
nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus.
Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

```
1  reverse :: SArr Int32 → Software ()
2  reverse arr =
3    do for 0 (len `div` 2) $ λix →
4         do aix ← getArr arr ix
5            ajx ← getArr arr (len - ix)
6            setArr arr ix          ajx
7            setArr arr (len - ix) aix
8    where
9      len = length arr
```

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetuer adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

```
1  reverse :: HArr Int32 → Hardware ()
```

Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper, leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellentesque, augue quis sagittis posuere, turpis lacus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla facilisi. Sed a turpis eu lacus commodo facilisis. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetuer.

```
1  reverse :: (Arrays m, Expr (Exp m), Type' (Exp m) Int32) ⇒
2    Arr m (Exp m Int32) → m ()
```

Suspendisse vel felis. Ut lorem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, molestie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque vehicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.

Chapter 3

Concluding Remarks

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetuer id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetuer adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae

tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper, leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellentesque, augue quis sagittis posuere, turpis lacus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla facilisi. Sed a turpis eu lacus commodo facilisis. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetuer.

Suspendisse vel felis. Ut lorem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, molestie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque vehicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.

Sed commodo posuere pede. Mauris ut est. Ut quis purus. Sed ac odio. Sed vehicula hendrerit sem. Duis non odio. Morbi ut dui. Sed accumsan risus eget odio. In hac habitasse platea dictumst. Pellentesque non elit. Fusce sed justo eu urna porta tincidunt. Mauris felis odio, sollicitudin sed, volutpat a, ornare ac, erat. Morbi quis dolor. Donec pellentesque, erat ac sagittis semper, nunc dui lobortis purus, quis congue purus metus ultricies tellus. Proin et quam. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Praesent sapien turpis, fermentum vel, eleifend faucibus, vehicula eu, lacus.

Bibliography

Mudge, Trevor (2001). "Power: A first-class architectural design constraint". In: *Computer* 34.4, pp. 52–58 (cit. on p. 3).

Networking Index, Cisco Visual (2016). "Forecast and methodology, 2016-2021, white paper". In: *San Jose, CA, USA* (cit. on p. 3).

Obile, W (2016). *Ericsson Mobility Report* (cit. on p. 3).

Persson, Anders (2014). "Towards a functional programming language for baseband signal processing". In: (cit. on p. 3).

Part II

Appended papers

Paper 1

Stream Processing for Embedded Domain Specific Languages

Markus Aronsson, Emil Axelsson, Mary Sheeran

Stream Processing for Embedded Domain Specific Languages

Markus Aronsson   Emil Axelsson   Mary Sheeran

Chalmers University of Technology
mararon@chalmers.se, emax@chalmers.se, ms@chalmers.se

Abstract

We present a library for expressing digital signal processing (DSP) algorithms using a deeply embedded domain-specific language (EDSL) in Haskell. The library supports definitions in functional programming style, reducing the gap between the mathematical description of streaming algorithms and their implementation. The deep embedding makes it possible to generate efficient C code. The signal processing library is intended to be an extension of the Feldspar EDSL which, until now, has had a rather low-level interface for dealing with synchronous streams. However, the presented library is independent of the underlying expression language, and can be used to extend any pure EDSL for which a C code generator exists with efficient stream processing capabilities. The library is evaluated using example implementations of common DSP algorithms and the generated code is compared to its handwritten counterpart.

Categories and Subject Descriptors   D.3.2 [*Language Classification*]: Applicative (functional) languages—Dataflow languages; C.3 [*Computer Systems Organization*]: Signal processing systems

General Terms   Languages, Design

Keywords   Digital Signal Processing, Stream processing, Embedded domain specific language, dataflow, synchronous programming, observable sharing, code generation

1.   Introduction

In recent years, the amount of traffic passing through the global communications infrastructure has been increasing at a rapid pace. Worldwide, total Internet traffic is estimated to grow at an average rate of 32% annually, reaching approximately eighty million terabytes per month by the end of 2015 [20]. For telecommunications infrastructure, the consequence of such a rapid growth rate has been a dramatic increase in the demand for network capacity and computational power [1].

Today, digital signal processing software is typically implemented in low level C, which forces designers to focus on low-level implementation details rather than the mathematical specification of the algorithms. Our group is developing an embedded domain-specific language, Feldspar [5], that aims to raise the abstraction level of signal processing software by expressing algorithms as pure functional programs.

However, signal processing is more than just pure computations – it is also about how to connect those functions in a network that operates on streaming data. A suitable programming model for reactive systems that process streams of data is *synchronous dataflow* (SDF) [21], which offers natural, high-level descriptions of streaming algorithms, while still permitting the generation of efficient code. There already exist a number of SDF languages, such as Lustre [10], but we are interested in adding an SDF programming model to Feldspar rather than using an existing language. The reason is that Feldspar is an ecosystem of different programming models which together offer much more flexibility than a typical SDF language.

Feldspar does have a library for programming with synchronous streams, but they are implemented in such a way that we cannot analyze them or detect sharing – a common problem in embedded languages. As a result of this, we cannot prevent code duplication and instead need to resort to using low-level combinators, which are tedious to use and do not scale well. Our solution is not to abandon such streams altogether – instead, we wrap them in a layer that remedies these problems.

This paper describes a library for extending an existing Haskell EDSL, like Feldspar, with support for SDF. The underlying EDSL is used to represent pure functions (which, of course, can be arbitrarily complicated), and our library gives a means to connect such functions using an SDF programming model. If the underlying EDSL provides a C code generator with a given interface, our library is capable of emitting C code for SDF programs. While we are interested in using Feldspar as the expression language, the library is not dependent on Feldspar, and so may be of interest to other EDSL developers.

This paper makes the following contributions:

- We present an EDSL for synchronous dataflow programming in Haskell (Section 2). By allowing the dataflow nodes to run arbitrary code, our EDSL greatly increased the expressive power compared to previous dataflow EDSLs in Haskell, such as Lava [11, 16]. Practically, the result is a useful addition to Feldspar.

- We show how to abstract away from the underlying expression language by establishing an interface for the un-

derlying expression compiler and interpreter (Sections 3 and 4.6). As a result, we ensure that our EDSL and the expression language are loosely coupled, giving a clear separation of concerns and an interchangeable expression language.

- We show a method for code generation from our EDSL via two lower-level representations: streams and imperative programs (Section 4). Our representation of streams breaks away from the classical one [9], as streams instead take on a monadic form using imperative programs. This allows us to generate efficient code from streams.

Section 5 gives an evaluation of the method on some simple examples.

2. Signals

Our library[1] is based on the concept of *signals*: possibly infinite sequences of values in some pure expression language. The notion of *time* used is discrete and non-negative, as we require signals to be causal: output may only depend on current or previous input. If the output of a signal, or signal function, does not depend on previous input, it is called *combinatorial*. Otherwise it is *sequential*.

Conceptually, signals can be thought of as infinite lists, similar to those found in most lazy functional languages:

```
Sig a  ≈  [a]
```

Unlike lists, `Sig` is, however, not a first-class value, as nesting of signals is disallowed.

Programming of signals is done compositionally; a rich set of operators supports the composition of new signals from existing ones. Signal programs are a collection of mutually recursive signal functions, each built from static values or other signals. For instance, some of the supplied combinatorial functions include:[2]

```
repeat :: Typeable a ⇒ Expr a → Sig a

map :: (Typeable a, Typeable b)
    ⇒ (Expr a → Expr b)
    → Sig a → Sig b

zipWith
    :: (Typeable a, Typeable b, Typeable c)
    ⇒ (Expr a → Expr b → Expr c)
    → Sig a → Sig b → Sig c
```

where `repeat` constructs a signal by repeating some value, `map` promotes a function to operate element-wise over signals, and `zipWith` joins two signals. These functions represent different kinds of nodes in the signal graph.

A fixed expression language, `Expr`, has been used in the above interface in order to improve readability. Later on, in section 4.6, this interface will be generalized to arbitrary expression types. Also, as a result of using type casting internally, values are required to be `Typeable`. Such a constraint is not harmful in practice, as most types satisfy the constraint.

We make use of existing type classes where possible, as they provide a familiar user interface when composing signals. For instance, numerical operations can be supported for signals by instantiating Haskell's `Num` class:

---

[1] Our library is available as open source [3].

[2] Because we are reusing names from Haskell's `Prelude`, these names must be hidden when using our library. This can be done by beginning each program with `import qualified Prelude`.

```
instance (Num a, Typeable a) ⇒ Num (Sig a)
  where
    fromInteger = repeat . fromInteger
    (+)         = zipWith (+)
    (-)         = zipWith (-)
    ...
```

The general idea is that every ground type, as permitted by the expression language, is lifted to operate element-wise over signals. This can either be done by redefining standard Haskell functions, as with the previous `repeat` and `map` functions, or through type classes.

For classes that cannot be instantiated in this way, we instead provide custom classes that override the default ones. For instance, as the `Eq` class has methods with non-overloaded types, making it incompatible with `Sig`, a customized version of that class is provided:

```
class EEq a
  where
    (==:) :: Expr a → Expr a → Expr Bool
    (/=:) :: Expr a → Expr a → Expr Bool
```

Using standard classes and function names in this way simplifies the construction of signals by providing a homogeneous user interface; complex networks can be defined using standard Haskell functions.

So far, only combinatorial functions have been considered. Most interesting examples of signal programs do, however, carry some form of state. For this reason, a sequential operator is provided:

```
delay :: Typeable a
      ⇒ Expr a → Sig a → Sig a
```

`delay` prepends a value to a signal, effectively delaying the input signal by one instant. Note that `delay` introduces the notion of a *next time step*, making time enumerable.

While it may appear innocent, the combination of `delay` with feedback allows one to express any kind of sequential network. For example, an edge detector could be implemented in terms of the delay function, checking at each instant if the values have changed:

```
edge :: Sig Bool → Sig Bool
edge s = zipWith (/=) s (delay false s)
```

where `false` creates a boolean value in the expression language.

An edge detector is however quite a small example. Therefore, in order to develop a better understanding of how programming with the signal library works, we will consider some larger examples, including the use of feedback.

2.1 FIR Filter

Finite impulse response (FIR) filters are one of the two primary types of digital filters used in DSP applications [23].

Consider the mathematical definition of a FIR filter of rank N:

$$y_n = \sum_{i=0}^{N} b_i \cdot x_{n-i} \qquad (1)$$

While this description is already convenient for software realization, representing the decomposition graphically can yield a better intuition of what constitutes the filter.

The resulting filter is depicted in Figure 1, where the z^{-1} block is a unit delay. Looking at this representation, where
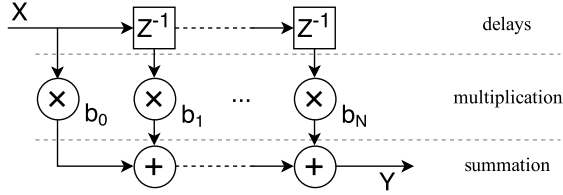
Figure 1. A direct form discrete-time FIR filter of order N

each column represents a term in the expanded summation, it is clear that the filter consists of three components: a number of successive unit delays, multiplication with coefficients and a summation.

Numerical operations are already supported for signals, since they instantiate Haskell's `Num` class. Consequently, in order to implement a FIR filter for signals, we only require a couple of helper functions to model each of its main components. These components are simply repeated applications of existing signal functions, and can therefore be modeled using pure Haskell functions:

```
import qualified Prelude as P


sums :: (Num a, Typeable a) ⇒ [Sig a] → Sig a
sums = P.foldr1 (+)


muls :: (Num a, Typeable a)
     ⇒ [Expr a] → [Sig a] → [Sig a]
muls as = P.zipWith (*) (P.map repeat as)


dels :: Typeable a ⇒ Expr a → Sig a → [Sig a]
dels a = P.iterate (delay a)
```

Here, addition and multiplication are overloaded operators from the `Num` class, and functions prefixed with `P` are standard library functions in Haskell.

Given these functions, a FIR filter can be expressed as:

```
fir :: (Num a, Typeable a)
    ⇒ [Expr a] → Sig a → Sig a
fir as = sums . muls as . dels 0
```

which is quite close to the filter's graphical representation. Domain experts in digital signal processing tend to be comfortable with composing sub-components in this way.

A FIR filter is still a small example, and, more importantly, it lacks feedback, a necessary component in many signal processing applications. We will therefore present one more example, which includes the use of feedback.

2.2 IIR Filter

Infinite impulse response (IIR) filters comprise the second primary type of digital filters used in DSP, and, unlike FIR filters, they contain feedback.

An IIR filter is typically described and implemented in terms of a difference equation, which defines how the output signal is related to the input signal:

$$y_n = \frac{1}{a_0} \cdot \left(\sum_{i=0}^{P} b_i \cdot x_{n-i} - \sum_{j=1}^{Q} a_j \cdot y_{n-j} \right) \qquad (2)$$

P and Q are the feedforward and feedback filter orders, respectively, and a_j and b_i are the filter coefficients. Neverthe-

less, representing the filter graphically can once again yield a better understanding of what constitutes the filter.
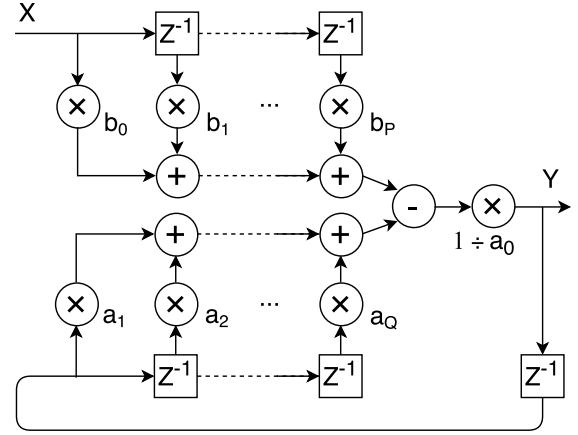

Figure 2. A direct form discrete-time IIR filter of order P and Q

Examining Figure 2, we see that this IIR filter loosely consists of two FIR filters, where the bottom one has an extra delay and is recursively defined. Also, the filter contains an extra subtraction and multiplication. This leads directly to the following code:

```
iir :: (Fractional a, Typeable a)
    ⇒ Expr a → [Expr a] → [Expr a]
    → Sig a → Sig a
iir a0 as bs x = y
  where
    y = (1 / repeat a0) * (upper x - lower y)
    upper = fir bs
    lower = fir as . delay 0
```

Circular definitions like this are possible because of the `delay` operator which makes sure that each output only depends on previous inputs, thus ensuring a productive network. In contrast, when defining our helper functions, `sums`, `muls` and `dels`, a different kind of recursion was used, leading to unfolding of the signal graph instead of feedback.

In general we have that recursively defined signals introduce feedback, while recursion over other Haskell values can be used to build graphs instead.

3. Embedding Programs

Before going into the implementation of signals, we will show how to embed an imperative programming language in Haskell. The choice of an imperative model is largely based on our interest in generating C code, but it could just as well have been some other sequential model. This embedded language will then form the basis for our implementation of signals.

The imperative language is implemented using the Operational package [2], which provides a monad, `Program cmd`, parameterized on its primitive instructions `cmd`. The idea behind Operational is that monadic programs can be viewed as sequences of instructions to be executed by some machine, so the task of implementing a monad is then the same as writing an interpreter for its instructions. It is possible to give multiple interpretations of the same program, and in

particular, it is possible to implement code generation just as a special program interpreter.

Our approach to code generation is similar to previously published methods [26, 27]. In these, code generation is implemented as a recursive function over a monadic program's graph representation, where each node corresponds to a monadic bind operation. Although Operational provides a similar functionality, we have chosen a different, and in some respects simpler approach – namely to use the following function:

```
interpretWithMonad :: Monad m
                   ⇒ (∀ a.      cmd a → m a)
                   → (Program cmd b → m b)
```

This function lifts a monadic interpretation of the primitive instructions, which may be of varying types, to a monadic interpretation of the whole program. By using different types for the monad, m, one can implement different "back ends" for the programs. For example, interpretation in the IO monad gives a way to run the programs in Haskell, and interpretation in a code generation monad can be used to make a compiler to another language.

We will demonstrate our approach to code generation using a simple example: reading and writing numbers through IO. First, we define a concrete expression language to use:

```
data Expr a where
  Var :: String → Expr a
  Lit :: Integral a ⇒ a → Expr a
  Add :: Integral a ⇒ Expr a → Expr a
                                → Expr a
  deriving (Typeable)
```

These constructors represent numeric literals, variables and additions. Expr is indexed on the result type of the expression, and by using a generalized algebraic data type (GADT) we make sure that only well-typed expressions can be represented (barring the fact that variables can have any type, independent of the environment). As the instruction set, the following type is used:

```
data CMD a where
  Get :: CMD (Expr Int)
  Put :: Expr Int → CMD ()
```

where CMD is an instruction type with two commands: Get for reading from stdin and Put for writing to stdout.

By using CMD as the parameter to Program, a specialized monad for the imperative EDSL is made:

```
type Prog a = Program CMD a
```

which also hides the instruction set used from the end user.

For simple code generation to C, one needs a combination of the State and Writer monads, the former to thread a fresh name supply and the latter to collect a sequence of statements; C statements are represented as strings for simplicity:

```
type Code = StateT Int (Writer [String])
```

Two helper functions are also needed, one for generating fresh names and another for compiling pure expressions; again, C expressions are represented as strings:

```
gensym :: Code String
gensym = do v ← get; put (v + 1);
              return ("v" ++ show v)
```

```
compExpr :: Expr a → String
compExpr (Lit a) = show $ toInteger a
compExpr (Var v) = v
compExpr (Add a b)
    = "(" ++ compExpr a ++ " + "
          ++ compExpr b ++ ")"
```

Compilation of primitive instructions is then given by the functional compCMD:

```
compCMD :: CMD a → Code a
compCMD Get = do
    v ← gensym
    tell [v ++ " = getchar();"]
    return (Var v)
compCMD (Put a) =
    tell ["putchar(" ++ compExpr a ++ ");"]
```

The compCMD function generates code for a single instruction, and is lifted for compilation of whole programs:

```
compile' :: Prog a → Code a
compile' = interpretWithMonad compCMD


compile :: Prog a → String
compile = unlines . execWriter
        . flip execStateT 0 . compile'
```

As an example of what the generated code looks like, consider the following small program:

```
gett = singleton Get    :: Prog (Expr Int)
putt = singleton . Put :: Expr Int → Prog ()


prg :: Prog ()
prg = do a ← gett; b ← gett;
         putt (a + b + 1)
```

It gets two numbers and then puts their sum plus one, where gett and putt use of Operational's singleton function to lift their instructions into actions. The expression a + b + 1 makes use of the following instance:

```
instance Integral a ⇒ Num (Expr a) where
    fromInteger = Lit . fromInteger
    (+)         = Add
    ...
```

The code generated from prg looks as follows:

```
*Main> putStr $ compile prg
v0 = getchar();
v1 = getchar();
putchar(((v0 + v1) + 1));
```

3.1 Abstracting away from the Expression Language

Note that the only thing we need to know about the Expr type in order to implement compCMD is that it has a Var constructor and a compilation function compExpr. So, to abstract away from the Expr type, we simply put those functions behind an abstract interface:

```
class CompExpr exp where
  type VarPred exp :: * → Constraint
  varExpr  :: VarPred exp a ⇒ String → exp a
  compExpr :: exp a → String
```

To account for the fact that different languages may have different constraints on the variable constructor, the asso-

ciated `VarPred` type allows us to use different constraints in each instance.[3] For example, in the instance for `Expr`, `VarPred Expr` maps to `Integral` due to the `Integral` constraint on the `Var` constructor:

```
instance CompExpr Expr where
  type VarPred Expr = Integral
  varExpr         = Var
  -- compExpr as before
```

We can now parameterize on the expression type used in `CMD`, in order to show that the technique described in this paper is independent of the expression language used.

```
data CMD exp a where
  Get :: CMD exp (exp Int)
  Put :: Expr Int → CMD exp ()
```

where the first parameter to `CMD` is the representation of pure expressions. Reimplementing `compCMD` using the abstract interface yields a compilation function with the following type:

```
compile
    :: (CompExpr exp, VarPred exp Int)
  ⇒ Program (CMD exp) a → String
```

As the type says, this version of `compileProg` is capable of compiling programs over *any* expression language, as long as it implements the `CompExpr` interface. Note the constraint `VarPred exp Int`, coming from the fact that `compCMD` needs to return a variable of type `exp Int` in the case for `Get`.

3.2 Running Programs

As previously noted, it is possible to run embedded programs in Haskell by providing an interpretation in the `IO` monad. Naturally, we would prefer to run such programs without constraining the expression language unnecessarily. A type class reminiscent of `CompExpr` is therefore introduced:

```
class EvalExpr exp where
  type LitPred exp :: * → Constraint
  litExpr  :: LitPred exp a ⇒ a → exp a
  evalExpr :: exp a → a
```

An instance for `Expr` is also made[4]:

```
instance EvalExpr Expr where
  type LitPred Expr  = Integral
  litExpr            = Lit
  evalExpr (Lit a)   = a
  evalExpr (Add a b) = evalExpr a
                     + evalExpr b
```

Based on `EvalExpr`, one can now conveniently define the behaviour of each primitive instruction and then lift that to get a run function for programs:

```
runCMD :: (EvalExpr exp, LitPred exp Int)
       ⇒ CMD exp a → IO a
runCMD Get     = fmap (litExpr . fromEnum)
                 $ getChar
runCMD (Put a) = putChar $ toEnum
                 $ evalExpr a
```

---

[3] `VarPred` requires the `ConstraintKinds` extension in GHC.

[4] Note that we assume that evaluated expressions do not contain any variables. That is true for the code in this paper, as variables are only introduced during code generation, but there is nothing preventing us from violating this assumption in general.

```
runProg :: (EvalExpr exp, LitPred exp Int)
        ⇒ Program (CMD exp) a → IO a
runProg = interpretWithMonad runCMD
```

Now that we have an embedding of imperative programs, we move on to describe the implementation of the signal library. By building it on top of the imperative embedding, we essentially get code generation for free.

4. Implementation

This section presents how signals are implemented and compiled to C code. While we stated earlier that signals represent sequences of values, the underlying implementation is a bit more involved: a signal consists of three different layers. At the bottom of these layers is the program layer, described in section 3. On top of that is the stream layer, a shallow implementation of co-iterative streams. Then, at the very top, is the signal layer, which extends the stream layer with support for observable sharing and rewriting.

4.1 Program Layer

Although our implementation is based on the same principle as the example in Section 3, it uses a more powerful expression language with support for many numerical and logical operations. The instruction set, `CMD`, has been extended as well and contains many more instructions – including ones with nested programs, such as this if statement whose branches are themselves programs:

```
data CMD exp a where
  ...
  If :: exp Bool
     → Program (CMD exp) a  -- true branch
     → Program (CMD exp) a  -- false branch
     → CMD exp a
```

While these extra instructions are useful for constructing more interesting examples, signals themselves only require basic functionality from the embedded language. More specifically, signals make use of mutable references – similar to `IORef` in Haskell. The interface to mutable references is given by the following functions:

```
newRef :: Typeable a
       ⇒ Expr a → Prog (Ref (Expr a))
getRef :: Typeable a
       ⇒ Ref (Expr a) → Prog (Expr a)
setRef :: Typeable a
       ⇒ Ref (Expr a) → Expr a → Prog ()
```

where `newRef` creates a new reference, `getRef` reads from a reference and `setRef` writes to one.

4.2 Stream Layer

The stream layer is an implementation of co-iterative streams, as described, for instance, by Caspi and Pouzet [9]. Co-iteration consists of associating to each stream a transition function from old state to a pair of a value and a new state, and an initial state. Thus, in Haskell we can model co-iterative streams as follows:

```
data Stream a where
    Stream :: (state → (a, state))
           → state
           → Stream a
```

This approach allows us to handle infinite data types, like streams, in a strict and efficient way, instead of having to deal with them using lazy data structures such as lists. As a result of this, stream transformers are free to pick apart the input streams in any way they see fit. Fusion of streams is achieved by construction.

A slightly modified version of these co-iterative streams is used in our library, where the state is represented implicitly by the `Program` monad:

```
data Stream a = Stream (Prog (Prog a))
```

In this definition, the outer monad is used to initialize the stream, and the result of the initialization is a monadic action that can be executed repeatedly to produce the stream of values. Since the program layer supports stateful computations, this representation allows for the state to be updated while running. Since we often have streams with `Expr` in the result type, we define a convenient shorthand:

```
type Str a = Stream (Expr a)
```

We can use this type to define stream versions of the `repeat` functions from Section 2:

```
repeatStr :: Expr a → Str a
repeatStr a = Stream $ return $ return a
```

Since `repeatStr` does not carry any state, it has an empty initialization and always produces the same value, `a`. In a similar manner, we can also define a stream version of the `map` function:

```
mapStr :: (Expr a → Expr b) → Str a → Str b
mapStr f (Stream init) = Stream $ do
  next ← init
  return $ do
    a ← next
    return (f a)
```

Unlike `repeatStr`, `mapStr` carries state, or at least the state of its argument. As a result, `init` is run to initialize the stream and access the argument's transition function `next`. Output is then produced by running `next`, transforming each value it produces using `f` before returning it.

These are the combinatorial functions we can implement using streams, but as the `Program` type also supports mutable state, it is also possible to have sequential, state-carrying, nodes. For instance, we can implement a function corresponding to `delay` from Section 2:

```
delayStr :: Typeable a ⇒
            Expr a → Str a → Str a
delayStr a (Stream init) = Stream $ do
  next ← init
  r    ← newRef a
  return $ do
    v ← next
    b ← getRef r
    setRef r v
    return b
```

This definition is similar to `mapStr`, but instead of transforming values, it stores them to be returned in the next iteration. Note how the local state is created in the initialization stage, after intializing the input stream.

By defining a `zipWithStr` fucntion, analogously to `mapStr`, we can reimplement all functions from Section 2 for streams. However, this leads to problems: because streams are implemented as a shallow embedding on top of programs, we lose the ability to analyze and optimize the dataflow networks. The problem can be seen, for example, in the definition of `mapStr`, where the code for the output stream includes the whole code for the input stream. So the result of `mapStr` is code where we can no longer see where nodes begin and end.

This problem is more severe than it may seem at first:

- Nodes that are shared in the network will lead to duplication in the generated code.

- Related to the above, feedback loops will lead to code being duplicated indefinitely.

Additionally, there are many optimizations that can be performed if the network can be analyzed. For example, it is more efficient to implement delay lines, as generated by the earlier `dels` function, using circular buffers instead of individual delay elements. We are interested in doing such rewriting optimizations automatically.

4.2.1 Feldspar's Solution

It is possible to work around the above problems by defining combinators for creating specific types of networks. This is what was done earlier in Feldspar. Feldspar has a stream library[5] implemented roughly as described in this section.

One of Feldspar's network combinators is `recurrenceIO`, defining a feedback network with one input and one output:

```
recurrenceIO :: (...)
    ⇒ Vector1 a        -- initial in buffer
    → Stream (Data a)  -- input stream
    → Vector1 b        -- initial out buffer
    → (Vector1 a → Vector1 b → Data b)
    → Stream (Data b)  -- output stream
```

The fourth argument represents the actual computation in the network: a function that, given vectors of previous inputs and outputs, computes the next output. The sizes of the initial buffers determine the number of previous values that are remembered. Restricting access to previous values in this way enables efficient memory management, and the vectors are implemented as circular buffers internally.

Given this `recurrenceIO` function, it is possible to express recursively defined streams. For instance, Feldspar's stream library includes the following definition of an IIR filter:

```
iir :: Data Float → Vector1 Float
    → Vector1 Float → Stream (Data Float)
    → Stream (Data Float)
iir a0 as bs x =
  recurrenceIO (replicate (length bs) 0) x
               (replicate (length as) 0)
    (λi o → 1 / a0 * ( scalarProd bs i
                       - scalarProd as o))
```

where `scalarProd` is the scalar product on Feldspar's vectors.

While functions such as `recurrenceIO` allow efficient networks to be expressed, they do not provide a very satisfactory solution to programming with streams. First of all, they are quite unintuitive to work with. Worse still, each combinator only captures a specific kind of network. For this reason, Feldspar provides a number of recurrence combinators differing only in the number of input and output streams they handle. All in all, the approach does not scale very well.

However, the solution to the above problems is not to abandon streams altogether – instead, we wrap streams in

---
[5] http://hackage.haskell.org/package/feldspar-language-0.7/docs/Feldspar-Stream.html

a layer using a deep embedding. This new layer will then remedy the problems inherent in a shallow model while still having access to streams and the nice properties they have for code generation.

4.3 Signal Layer

The last layer, called the signal layer, will act as our wrapper layer for the underlying stream model. As such, it provides a way to promote streams and stream functions into a corresponding signal version:

```
data Signal a where
  Const :: Typeable a
         ⇒ Str a → Signal (Expr a)

  Lift  :: (Typeable a, Typeable b)
         ⇒ (Str a → Str b)
         → (Signal (Expr a) → Signal (Expr b))


type Sig a = Signal (Expr a)
```

`Const` is used for lifting constant streams, while `Lift` lifts a stream transformer into a signal transformer. Like `Str`, `Sig` defines a shorthand for signals that return expressions.

Most of the signal library's functionality comes from lifting stream functions; for instance, the previous `map` and `repeat` functions are both implemented in this way:

```
repeat :: Typeable a ⇒ Expr a → Sig a
repeat a = Const (repeatStr a)

map :: (Typeable a, Typeable b)
    ⇒ (Expr a → Expr b) → Sig a → Sig b
map f = Lift (mapStr f)
```

Defining signal functions in this way introduces a ambiguity to signals: how to distinguish between pairs of signals and signals of pairs. Additional constructors are therefore included, for managing the flow of a signal and helping to differentiate between the two kinds of signals:

```
data Signal a where
  ...
  Zip :: (Typeable a, Typeable b)
       ⇒ Signal a → Signal b → Signal (a, b)
  Fst :: Typeable a
       ⇒ Signal (a, b) → Signal a
  Snd :: Typeable b
       ⇒ Signal (a, b) → Signal b
```

Manual zipping and unzipping can however get tiresome. Signals are therefore extended with a generalized constructor, `Map`, which allows one to have signal functions which take and produce arbitrary trees of signals. These trees are represented using the `Struct` type.

```
data Signal a where
  ...
  Map :: (Typeable a, Typeable b)
       ⇒ (Struct a → Struct b)
       → (Signal a → Signal b)


data Struct a where
  Leaf :: Typeable a
       ⇒ Expr a → Struct (Expr a)
  Pair :: Struct a → Struct b
       → Struct (a, b)
```

Armed with these two, it is possible to implement a more general zipping function `zipWith`, which joins the two signals element-wise using the given function:

```
zipWith
    :: (Typeable a, Typeable b, Typeable c)
    ⇒ (Expr a → Expr b → Expr c)
    → Sig a → Sig b → Sig c
zipWith f a b =
    Map (λ(Pair (Leaf a) (Leaf b))
              → Leaf (f a b)) $ Zip a b
```

Defining functions in this way is not as elegant as one would like it to be. It is possible, however, to automate this process, and rewrite the previous definition of `zipWith`:

```
zipWith f = curry $ lift $ uncurry f
```

What the `lift` does is to transform Haskell tuples into signals, using `Zip`, apply the function, and then transform them back again.

Sequential networks, as shown in section 2.2, can be expressed using recurrence equations. However, due to their limitations, we prefer to avoid using them entirely and instead program feedback using the more general `delay` function shown in section 2. A dedicated `Delay` constructor is therefore included in the `Signal` type:

```
data Signal a where
  ...
  Delay :: Typeable a
         ⇒ a → Signal a → Signal a


delay :: Typeable a
       ⇒ Expr a → Sig a → Sig a
delay = Delay
```

Explicitly describing feedback networks that include delays is much more user friendly than imposing a fixed set of combinators for recurrence equations.

4.4 Running signals

Now that we have defined all the layers, we would like to be able to peel them off. As signals are a wrapper for the underlying stream model, we want them, and the overhead associated with them, to disappear entirely during the code generation process.

This brings us up against a common problem in manipulating data types representing embedded languages: how to observe sharing and cycles in trees generated from a deeply embedded language. Programs in the signal library represent data flow graphs, which can contain shared nodes and cycles. But the `Signal` type is a tree. In order to view the `Signal` type as a graph, we need to examine how this tree is represented in memory (because even trees are represented as graphs internally by GHC).

One approach to observable sharing, as proposed by Gill [15], relies on GHC-specifics to provide an `IO` function capable of observing sharing directly. Given such a function, it is possible to define a function that strips away the signal layer and gives back a stream function. We call this the `compile` function:

```
compile :: (Typeable a, Typeable b)
        ⇒     (Sig a → Sig b)
        → IO (Str a → Str b)
compile f = do
  (Graph nodes root) ← reifyGraph f
```

```
    let links = linker nodes
        order = sorter root nodes
        cycle = cycles root nodes

    return $ if cycle
        then error "Cycle found"
        else compiler nodes links order
```

The function `reifyGraph` is the one that turns a signal function into its representation as a directed acyclic graph. This process is commonly known as reification: to take something abstract and regard it as material instead; a reified type is simply a value that represents a type. Reification of signals through observable sharing therefore implies that a recursive data type is used to represent the embedded language, like `Signal`, coupled with a mirror type where all points of recursion have been replaced by abstract references:

```
data TSignal r where
  TConst :: (Typeable a) ⇒ Str a → TSignal r
  TLift  :: (Typeable a, Typeable b)
          ⇒ (Str a → Str b) → r → TSignal r
  TZip :: r → r → TSignal r
  TFst :: r        → TSignal r
  ...

type Node = TSignal Int
```

Typing information is however lost after reification, as nodes are no longer connected by types. Typecasting is therefore used internally to connect these reified nodes, hence the need for `Typeable`.

As a means to recover some of the lost information, two accessory functions are used in the compiler: `linker` and `sorter`. As their names suggest, `linker` reconnects nodes with their expected input and output sources, while `sorter` sorts the nodes in topological order. Furthermore, as we cannot compile un-delayed feedback loops, `cycles` is used to check for valid signal networks before they are compiled.

Sorting and cycle detection of graphs are both well-studied problems and efficient solutions exist for both. Our `sorter` and `cycles` are implemented as follows: the graph is traversed in a depth-first manner, making sure to avoid loops by tagging nodes as they are visited. These two functions are given the following type signatures:

```
sorter :: Int → [(Int, Node)] → Map Int Int

cycles :: Int → [(Int, Node)] → Bool
```

Worth mentioning is the treatment of delays while detecting cycles, as cycles should only be rejected if they do not contain any delays. The solution we chose was to simply ignore all outgoing edges from a delay node. This breaks all "good" cycles and leaves the "bad" ones for us to detect. Also, as a result of edges being removed, we might end up splitting the graph into subgraphs but we can safely process these one at a time.

Lastly, we have the `linker` function, which reconnects nodes with their input sources, while also recovering parts of the typing lost during reification. More specifically, `linker` provides a typed `Struct` where each value has been replaced by a reference to the node that produces it.

Implementing `linker` as depth-first traversal of the graph – in the same manner as `sorter` or `cycles` – does however turn out to be unsatisfactory: `linker` must account for both the ordering of nodes and avoid cycles. As a result of this, its logic will be hidden behind all the necessary bookkeeping.

Instead, `linker` can be better expressed using a circular programming technique. We employ the same technique as Axelsson [4] to express circular knot-tying programs using a monad. This allows us to both declare connection constraints and read their solution at the same time, without having to worry about doing so in the correct order.

The general idea behind knot-tying is to make use of Haskell's laziness and a combination of the `Reader` and `Writer` monads:

```
import Control.Monad.Reader
import Control.Monad.Writer

type Knot r c m = ReaderT r (WriterT [c] m)
```

Constraints can now be declared through `Writer` and solutions are read using `Reader`. All that's left is to have Haskell's laziness tie the constraints and solution together:

```
tie :: MonadFix m ⇒ ([c] → r)
                    → Knot r c m a
                    → m (a, r)
tie solver knot = mdo
  (a, cs) ← runWriterT $ runReaderT knot r
  let r = solver cs
  return (a, r)
```

where `mdo` is notation for recursive monadic programs. The first argument to `tie` is the solver, which computes a solution from a list of constraints.

Applying the `Knot` monad to our problem of linking means we need to come up with a model for our constraints and solutions, which we would then write in bulk and have `tie` sort out their correct ordering. For this reason, a new type, called `TStruct`, is introduced to store the connection solutions. `TStruct` has a similar form to `Struct` but contains references at the leaves instead of values:

```
data TStruct a where
  TLeaf :: String → TStruct a
  TPair :: TStruct a → TStruct b
        → TStruct (a, b)
```

The type of these structs will however differ from node to node, which means that they cannot be stored directly in a plain map. Yet another datatype is therefore introduced, modeling existential types:

```
data EStruct where
  Ex :: Typeable a ⇒ TStruct a → EStruct

data Res = Resolution
  { _in  :: Map Int EStruct
  , _out :: Map Int EStruct }
```

Constraints are then simply members of these mappings, represented as:

```
data Cons = In (Int, EStruct) | Out (Int, EStruct)
```

As the knotty details of cycles and orderings are now entirely managed by knots, our `linker` can focus on describing the connections between nodes.

```
linker :: [(Int, Node)] → Res
linker = snd . runIdentity . tie solve
              . sequence    . map link
```

```
solve :: [Cons] → Res
solve cons = Resolution
  { _in  = fromList [i | In  i ← cons]
  , _out = fromList [o | Out o ← cons] }

link :: Monad m ⇒ (Int, Node)
                  → Knot Res Cons m ()
link (i,TConst  _) =
  tell [Out (i, Ex $ TLeaf $ show i)]
link (i,TLift _ _) =
  tell [Out (i, Ex $ TLeaf $ show i)]
link (i,TZip l r)  = do
  (Ex u) ← asks ((! l) . _out)
  (Ex v) ← asks ((! r) . _out)
  tell [Out (i, Ex $ TPair u v)]
...
```

Some type casting is required in `link`, but has been omitted to improve readability. Furthermore, nodes like `Zip`, which only forward their input references, can safely be removed from the final mapping as no other nodes will reference them – simplifying networks by rewriting them in a safe manner.

Linking nodes together in this manner works as long as we can reconstruct the expected input and output types. When the necessary information isn't available, we extend the constructor to include the missing pieces. For example, in order to link a `Map` node, we add two trees to its constructor:

```
data TSignal r where
  ...
  TMap :: (Typeable a, Typeable b)
       ⇒ TStruct a → TStruct b
       → -- same as before
```

Given these two pieces of extra information, we can implement `link` for `TMap` nodes:

```
...
link (i, TMap ti to f s) = do
  (Ex (u :: TStruct ti)) ← asks ((! s) . _out)
  tell [In  (i, Ex u)]
  tell [Out (i, Ex $ mark to $ show i)]

mark (TLeaf _)   s = TLeaf s
mark (TPair l r) s = TPair l' r'
  where l' = mark l $ s ++ "_l"
        r' = mark r $ s ++ "_r"
```

Compilation of streams is now straightforward, and our `compiler` function is given the following type:

```
compiler :: (Typeable a, Typeable b)
         ⇒ Map Int Node    -- nodes
         → Map Int EStruct  -- links
         → Map Int Int      -- order
         → (Str a → Str b)
```

The general idea is to simply traverse each node in the given order and lift their logic into programs as we go. This behavior is captured by the following pattern:

```
compileNode (id, node) = do
  -- fetch input addresses
  -- apply node logic
  -- store result in variables
  -- write said variables to output addresses
```

where input and output addresses are given by the links mapping. Also, since `linker` connected nodes directly with

their input sources, compiling pure flow nodes, like `Zip`, amounts to simply doing nothing at all:

```
compileNode (id, Zip {}) = return ()
```

Signal functions can now be compiled into streams using the previous `compile` function, which acts as an inverse of signal's `Lift`.

Then, as streams are no more than programs in disguise, we can generate code from them using the previously shown techniques. Inspecting the code produced in this way shows that all signal specific constructions, like delays, have indeed been removed by the compiler, leaving behind only arithmetic operations and control structures.

For example, applying the compiler to a rank-2 FIR filter produces the following code:

```
r6 = 0.0; r7 = 0.0; r3 = 0; r19 = 0;
while (r19 ≤ 999) {
    float a20;
    a20 = a0[r3];
    r8 = a20;
    r11 = 1.1;
    r10 = r11 * r8;
    r14 = -0.55;
    r15 = r6;
    r13 = r14 * r15;
    r17 = 0.275;
    r18 = r7;
    r16 = r17 * r18;
    r12 = r13 + r16;
    r9 = r10 + r12;
    r7 = r15;
    r6 = r8;
    a2[r3] = r9;
    r3 = r3 + 1;
    r19 = r19 + 1;
}
```

Here we only show the actual filter loop and omit all variable declarations. The input and output arrays are `a0` and `a2`, respectively. Variables `r14` and `r17` are the filter coefficients.

4.5 Internalized IO

So far, we have seen pure signal functions, such as `repeat`, `map` and `zipWith` and a single stateful function – `delay`. The pure functions are implemented using `Const` and `Lift`, and `delay` is given a specific constructor in the `Sig` type. However, there is in fact nothing stopping us from using `Const`/`Lift` to lift stateful or side-effecting functions to the signal layer. As an example, given a function that gets the current time of day as a number, we can easily make a signal source that gives the current time in each cycle:

```
getTime :: Prog (Expr Double)

time :: Sig Double
time = Const $ Stream $ return getTime
```

In a similar way, it would be possible to make signal functions for reading and writing to files or other resources. One complication then is that the lifted operations must be *non-blocking*, since the dataflow network must be able to keep running even when data is not available.

We have not yet explored the possibility of having networks with internalized IO, but it appears to be very convenient way of programming embedded systems. Rather than having our previous `compile` function with exactly one in-

put and one output signal, one could then instead have a function such as the following:

```
compile :: Typeable a
        ⇒ Sig a → IO (Program ())
```

The argument is a dataflow graph with all external communication built-in and the result is a monolithic program that can be invoked repeatedly to execute the dataflow network.

Although convenient for programming, the disadvantage of using internalized IO is that we can no longer view our networks as simple state machines with a pure transition function from previous to next state. With internalized IO, the "state" can easily reside outside of the program, and the transitions are generally non-deterministic (for example, the `time` source is unlikely to give the same sequence of numbers in two different runs of a system).

4.6 Abstracting away from the Expression Language

As we saw in Section 3.1, the underlying representation of imperative programs abstracts away from the representation of pure expressions. Similarly, while signals and streams have been hard-coded to use `Expr`, our actual implementation uses a more general type:

```
data Stream exp a =
      Stream (Program (CMD exp)
                  (Program (CMD exp) a))

data Signal exp a where
    Const :: Typeable a
             ⇒ Stream exp a → Signal exp a

    Lift  :: (Typeable a, Typeable b)
             ⇒ (Stream exp a → Stream exp b)
             → (Signal exp a → Signal exp b)
    ...
```

In the interface towards the user, we hide this generality and use types that are equivalent to the ones presented in this section:

```
type    Str a = Stream Expr (Expr a)
newtype Sig a = Sig (Signal Expr (Expr a))
```

Note that the code developed in this section is essentially the same when written with more general types.

5. Evaluation

The aim of this work is to be able to generate efficient code that can be used in applications with high performance demands. In order to evaluate the generated code, we have compared it to reference implementations written in C.[6] The referenced C code uses an efficient circular buffer to store previous input values. The IIR filter was implemented by ourselves by simply extending the FIR filter, hopefully without introducing accidental inefficiencies.

Each of our tests applies a filter to a 1000-element vector. In order to reduce the impact of system jitter, each test is repeated a large number of times (10000), and the time reported is the average over those iterations. The measurements were carried out on an Dell Latitude 7420 laptop with an Intel Core i7 processor using GCC version 4.9.1.

---

[6] The FIR filter's reference implementation was obtained from http://www.iowahills.com/A7ExampleCodePage.html
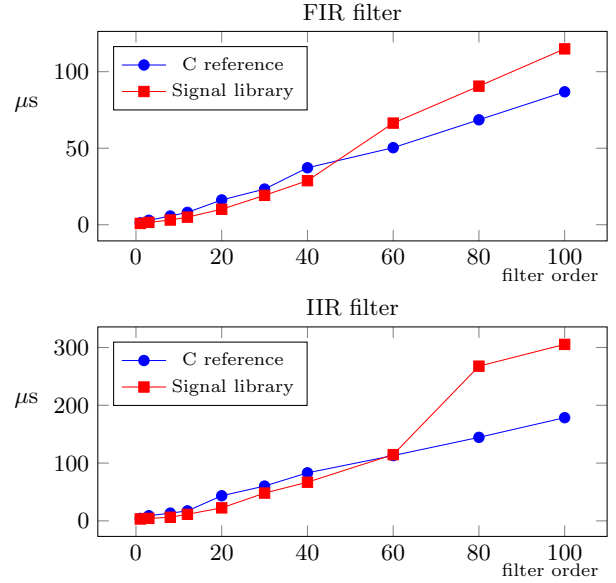

Figure 3. Running time of filters compared to reference C implementations.

The results for filters of varying orders are shown in Figure 3. The code from our library performs slightly better than the reference up to filter orders of around 50. For larger sizes it slows down because our library unrolls the inner loops of the filter, while the reference implementation uses nested loops. So the size of our code grows linearly with the filter order, and after a certain point the loop body becomes too big to be able to run efficiently (e.g. not enough registers to hold intermediate values).

The reason our library unrolls the inner loops is that we make use of recursive Haskell functions, such as `map` and `foldr1` to generate signal graphs. In order to generate nested loops instead, we would need to make use of more complex nodes, where, for example the multiply-accumulate part of the filter is represented as a single node whose input is an array. This in turn would require a more powerful expression language that has arrays and loops.

We are working on integrating the signal library with Feldspar. This will allow us to express filters with nested loops.

6. Related Work

As discussed in section 2, our library is based on the synchronous data flow paradigm and is heavily inspired by related languages from that paradigm, such as the Lucid Synchrone language [13, 24].

Lucid Synchrone is a member of the family of synchronous languages and is designed to model reactive systems. It was introduced as an extension of LUSTRE [17], and demonstrated that the language could be extended with new and powerful features. For instance, automatic clock and type inference were introduced, and a restricted form of higher-order functions was added.

However, Lucid Synchrone is a standalone language that cannot easily be integrated with EDSLs such as Feldspar. For this reason, we chose to implement a new library, partly

inspired by Lucid Synchrone, that brings an SDF programming model to existing EDSLs, such as Feldspar.

Another, rather different approach to modeling signal processing is used by Yampa [14]. Yampa is a member of the functional reactive programming paradigm, used for programming hybrid systems, that is, systems consisting of both continuous and discrete components. The key ideas in functional reactive programming are its notions of behaviors and events, where behaviors are continuous and time-varying, reactive values, and events are time-ordered sequences of discrete-time event occurrences [22].

Apart from the different notions of time, the behaviors of Yampa are quite similar to the synchronous concept of signals. Zélus [8], a successor of Lucid Synchrone, has shown that synchronous languages can indeed be extended to model hybrid systems as well. An event is usually a separate entity, modeling the control flow of a system. Some languages do, however, merge the two concepts, by describing events in terms of behaviors. This unification comes at a cost of some elegance, as behaviors can now be discrete, but it also means that it is possible to model such systems with only data flow.

Furthermore, functional reactive programming languages are commonly based on an applicative interface or arrows [19], which makes them a poor fit for a restricted EDSL like Feldspar. This comes as a result of type classes assuming that the underlying language is a super-set of Haskell, since every Haskell function can be promoted to an expression in the language.

Lava is a family of Haskell EDSLs designed for expressing hardware descriptions [7]. Chalmers Lava [11] is an experimental tool designed to assist circuit designers with hardware design, specification and verification.

Chalmers Lava is a structural hardware description language, embedded in Haskell in such a way that the designer actually writes and composes netlist generators. It is an archetypical deeply embedded language, with various backends operating on the data type representing circuits. Its most recent incarnation allows for the description of hardware by providing a polymorphic unit delay element and the ability to describe feedback (observable sharing). However, the other building blocks of ciruits are restricted to be simple gates. Here, in the signal library, we permit the combination of much more interesting building blocks, as our dataflow nodes can run arbitrary code – and indeed the users of the signal library can introduce their own program layer.

Kansas Lava [16] is another member of the Lava family of hardware description languages. It reimplements and extends the original Lava design patterns. One important extension is direct support for including new blocks of existing VHDL libraries as new, well typed primitives. Another noteworthy aspect of Kansas Lava is the use of a dual deep and shallow embedding (for synthesis and simulation) inside its `signal` type. This gives a clear design flow, starting from a Haskell function, adding applicative functors to give a model that understands time (shallow embedding) and then factoring out the applicative functors to give a synthesiable model (deep embedding).

Both of these Lava languages make use of observable sharing, where Chalmers Lava makes use of explicit references [12] and Kansas Lava uses Gill's type-safe version, which is also used in the signal library [15].

The ForSyDe (Formal System Design) methodology and tools [25] explore methods of developing heterogeneous systems in which different parts of the system may have differ-

ent models of computation (one of which is the synchronous model). An initial specification is refined into a detailed implementation model by the application of formally justified transformations. One of the implementations of ForSyDe is as an embedded language in Haskell and here, once again, the combination of deep and shallow implementations of a signal API is considered.

The way in which our program, stream and signal layers are layered one above the other is reminiscent of the Ivory and Tower EDSLs being developed at Galois Inc [18].

Ivory is an EDSL for generating safe, embedded C code, and it places particular emphasis on the definition of in-memory data structures and of how they should be manipulated. In our work on Feldspar, that aspect of C programming is covered largely by another "system layer" being developed to target a particular platform (the Adapteva Parallella board), above the Feldspar layer producing individual C functions. Ivory/Tower also has a layer for composing Ivory programs; the Tower layer covers the "glue code" that implements inter-process communication, processor locking, system clock reading etc. The utility and practicality of the layered EDSL approach is very well explained in the Galois experience report on this work [18].

In the work on the signal library described here, we have more and "thinner" layers than in the Ivory/Tower work. We expect to have yet more layers above the signal layer, although we have not yet decided exactly how to combine our various system layers.

7. Discussion

Writing digital signal processing software in C is a potentially tedious and error-prone task, primarily because of the enforced focus on low-level implementation details. We have described a library that alleviates this problem by extending an EDSL modeling pure functions with support for synchronous dataflow operations, allowing many low level details to be handled by the relevant compilers.

By using an EDSL for pure computations and Haskell's type classes, we can retain the elegance and modularity of traditional functional programming for signals. We are mainly interested in using Feldspar as the expression language; however, the library is not dependent on Feldspar, and so may be of interest to other EDSL developers.

We also intend to further build on earlier work on Lucid Synchrone (and LUSTRE) by introducing clocks as types [6]. This will allow us to deal correctly with sampling and oversampling, and to describe and implement multi-clocked networks. We feel that this will result in an interesting and expressive programming language for our envisioned users.

Having made the step to clocks, we are interested in the question of how to generate circuits rather than programs from our dataflow networks.

The intention is to grow the program layer to cover a substantial part of C. In order to support more complex systems, we also need a more powerful expression language. However, rather than extending the `Expr` type, our plan is to use the existing Feldspar EDSL [5] to represent pure expressions. This will require more advanced interleaving between the compilation phases for programs and expressions, as Feldspar expressions may generate complex chunks of C code.

Even if Feldspar ends up as our only expression language, we feel that the idea of establishing an interface towards the expression compiler (i.e. the `CompExpr` and `EvalExpr` classes

from Section 3.1) is a good one, as it allows us to treat the two parts of the compiler separately.

Acknowledgments

References

[1] Cisco visual networking index: Forecast and methodology, 2012–2017, May 2013. URL http://www.cisco.com/c/en/us/solutions/collateral/service-provider/ip-ngn-ip-next-generation-network/white_paper_c11-481360.pdf.

[2] H. Apfelmus. The operational package, version 0.2.3.2. URL https://hackage.haskell.org/package/operational.

[3] M. Aronsson. Signal processing for embedded domain specific languages, March 2015. URL https://hackage.haskell.org/package/signals.

[4] E. Axelsson. *Functional programming enabling flexible hardware design at low levels of abstraction*. PhD thesis, Chalmers University of Technology, 2008.

[5] E. Axelsson, K. Claessen, G. Dèvai, Z. Horvàth, K. Keijzer, B. Lyckegård, A. Persson, M. Sheeran, J. Svenningsson, and A. Vajda. Feldspar: A Domain Specific Language for Digital Signal Processing algorithms. In *Formal Methods and Models for Codesign (MEMOCODE), 2010 8th IEEE/ACM International Conference on*, pages 169–178, July 2010.

[6] D. Biernacki, J.-L. Colaco, G. Hamon, and M. Pouzet. Clock-directed Modular Code Generation of Synchronous Data-flow Languages. In *ACM International Conference on Languages, Compilers, and Tools for Embedded Systems (LCTES)*, June 2008.

[7] P. Bjesse, K. Claessen, M. Sheeran, and S. Singh. Lava: hardware design in Haskell. In *Proceedings of the third ACM SIGPLAN International Conference on Functional Programming*, ICFP '98, pages 174–184. ACM, 1998. ISBN 1-58113-024-4. URL http://doi.acm.org/10.1145/289423.289440.

[8] T. Bourke and M. Pouzet. Zélus: A Synchronous Language with ODEs. In *16th International Conference on Hybrid Systems: Computation and Control (HSCC'13)*, pages 113–118, Mar. 2013. URL http://www.di.ens.fr/~pouzet/bib/hscc13.pdf.

[9] P. Caspi and M. Pouzet. A co-iterative characterization of synchronous stream functions. *Electronic Notes in Theoretical Computer Science*, 11(0):1 – 21, 1998. ISSN 1571-0661. URL http://www.sciencedirect.com/science/article/pii/S1571066104000507.

[10] P. Caspi, D. Pilaud, N. Halbwachs, and J. A. Plaice. Lustre: A declarative language for real-time programming. In *Proceedings of the 14th ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages*, POPL '87, pages 178–188. ACM, 1987. ISBN 0-89791-215-2. URL http://doi.acm.org/10.1145/41625.41641.

[11] K. Claessen. *Embedded Languages for Describing and Verifying Hardware*. PhD thesis, Chalmers University of Technology, 2001.

[12] K. Claessen and D. Sands. Observable sharing for functional circuit description. In P. Thiagarajan and R. Yap, editors, *Advances in Computing Science — ASIAN'99*, volume 1742 of *Lecture Notes in Computer Science*, pages 62–73. Springer Berlin Heidelberg, 1999. ISBN 978-3-540-66856-5. URL http://dx.doi.org/10.1007/3-540-46674-6_7.

[13] J.-L. Colaço, A. Girault, G. Hamon, and M. Pouzet. Towards a Higher-order Synchronous Data-flow Language. In *Proceedings of the 4th ACM International Conference on Embedded Software*, EMSOFT '04, pages 230–239. ACM, 2004. ISBN 1-58113-860-1. URL http://doi.acm.org/10.1145/1017753.1017792.

[14] A. Courtney, H. Nilsson, and J. Peterson. The Yampa Arcade. In *Proceedings of the 2003 ACM SIGPLAN Workshop on Haskell*, Haskell '03, pages 7–18. ACM, 2003. ISBN 1-58113-758-3. URL http://doi.acm.org/10.1145/871895.871897.

[15] A. Gill. Type-safe Observable Sharing in Haskell. In *Proceedings of the second ACM SIGPLAN Symposium on Haskell*, Haskell '09, pages 117–128. ACM, 2009. ISBN 978-1-60558-508-6. URL http://doi.acm.org/10.1145/1596638.1596653.

[16] A. Gill, T. Bull, G. Kimmell, E. Perrins, E. Komp, and B. Werling. Introducing Kansas Lava. In M. Morazán and S.-B. Scholz, editors, *Implementation and Application of Functional Languages*, volume 6041 of *Lecture Notes in Computer Science*, pages 18–35. Springer Berlin Heidelberg, 2010. ISBN 978-3-642-16477-4. URL http://dx.doi.org/10.1007/978-3-642-16478-1_2.

[17] N. Halbwachs. *Synchronous Programming of Reactive Systems*. Number 215 in International Series in Engineering and Computer Science. Springer, 1993.

[18] P. C. Hickey, L. Pike, T. Elliott, J. Bielman, and J. Launchbury. Building embedded systems with embedded DSLs. In *Proceedings of the 19th ACM SIGPLAN International Conference on Functional Programming*, ICFP '14, pages 3–9. ACM, 2014. ISBN 978-1-4503-2873-9. URL http://doi.acm.org/10.1145/2628136.2628146.

[19] J. Hughes. Generalising Monads to Arrows. *Science of Computer Programming*, 37(1–3):67 – 111, 2000. ISSN 0167-6423. URL http://www.sciencedirect.com/science/article/pii/S0167642399000234.

[20] ITU. Global ITC development, 2001-2014, 2014. URL http://www.itu.int/en/ITU-D/Statistics/Documents/statistics/2014/stat_page_all_charts_2014.xls.

[21] E. Lee and D. Messerschmitt. Synchronous data flow. *Proceedings of the IEEE*, 75(9):1235–1245, Sept 1987. ISSN 0018-9219.

[22] H. Nilsson, A. Courtney, and J. Peterson. Functional Reactive Programming, Continued. In *Proceedings of the 2002 ACM SIGPLAN Workshop on Haskell*, Haskell '02, pages 51–64. ACM, 2002. ISBN 1-58113-605-6. URL http://doi.acm.org/10.1145/581690.581695.

[23] A. V. Oppenheim, R. W. Schafer, J. R. Buck, et al. *Discrete-time signal processing*, volume 2. Prentice-hall Englewood Cliffs, 1989.

[24] M. Pouzet. Lucid Synchrone, version 3. *Tutorial and reference manual. Université Paris-Sud, LRI*, 2006. URL http://www.di.ens.fr/~pouzet/lucid-synchrone/manual_html/.

[25] I. Sander. *System Modeling and Design Refinement in ForSyDe*. PhD thesis, Royal Institute of Technology, Stockholm, Sweden, April 2003. URL http://web.it.kth.se/~ingo/Papers/Thesis_Sander_2003.pdf.

[26] N. Sculthorpe, J. Bracker, G. Giorgidze, and A. Gill. The constrained-monad problem. In *Proceedings of the 18th ACM SIGPLAN International Conference on Functional Programming*, ICFP '13, pages 287–298. ACM, 2013. ISBN 978-1-4503-2326-0. URL http://doi.acm.org/10.1145/2500365.2500602.

[27] J. D. Svenningsson and B. J. Svensson. Simple and Compositional Reification of Monadic Embedded Languages. In *Proceedings of the 18th ACM SIGPLAN International Conference on Functional Programming*, ICFP '13, pages 299–304. ACM, 2013. ISBN 978-1-4503-2326-0. URL http://doi.acm.org/10.1145/2500365.2500611.

Paper 2

Hardware Software Co-design in Haskell

Markus Aronsson, Mary Sheeran

Hardware Software Co-design in Haskell

Markus Aronsson
Mary Sheeran
Chalmers University of Technology
Sweden

Abstract

We present a library in Haskell for programming Field Programmable Gate Arrays (FPGAs), including hardware software co-design. Code for software (in C) and hardware (in VHDL) is generated from a single program, along with the code to support communication between hardware and software. We present type-based techniques for the simultaneous implementation of more than one embedded domain specific language (EDSL). We build upon a generic representation of imperative programs that is loosely coupled to instruction and expression types, allowing the individual parts to be developed and improved separately. Code generation is implemented as a series of translations between progressively smaller, typed EDSLs, safeguarding against errors that arise in untyped translations. Initial case studies show promising performance.

CCS Concepts • **Software and its engineering** → **Functional languages**; **Domain specific languages**; • **Hardware** → *Reconfigurable logic and FPGAs*;

Keywords hardware software co-design, domain specific language

1 Introduction

Field Programmable Gate Arrays (FPGAs) represent an interesting point on the spectrum between general purpose processors and application specific integrated circuits (ASICs). An ASIC implements a fixed function, while an FPGA can be reprogrammed again and again, giving varying functionality over time. Such a programmable circuit typically consists of a large array of configurable logic blocks, connected by programmable interconnects. However, modern FPGAs also contain various discrete components, such as blocks of RAM, digital signal processing slices, and processor cores. Indeed, a modern FPGA can be viewed as a prototypical *heterogeneous system* that mixes components (such as processors and logic blocks) that need to be programmed in different ways.

FPGAs can combine the benefits of highly parallel, high performance hardware-based systems with the programmability of software, and at an attractive price point. In both embedded systems and high performance computing, FPGAs are increasingly combined with accelerators such as graphics processing units, and with multicores. And soon, there will be FPGAs, GPUs and CPUs all on one chip. The rise of FPGAs, while inexorable, has been slowed by the fact that they are difficult to program. The logic blocks are usually programmed in a hardware description language (HDL) such as VHDL or Verilog, while the embedded cores are typically programmed in a low level dialect of C. This mixture of hardware and software development is often called *co-design* (Teich 2012). The developer must specify both hardware and software parts, and how they communicate; ideally she would like to experiment with various choices of what to put in hardware and what in software. Current tools provide little support for such design exploration.

This paper presents first steps towards a system in which the entire design process, including exploration to decide where the boundary between hardware and software should be, is done in Haskell. We aim to give the user a reasonable degree of control over the generated code, both in hardware and software, by generating C and VHDL and making use of a standard Register Transfer Level (rather than high level) synthesis tool to produce the final hardware implementation. We build upon earlier work on code-generating embedded domain specific languages (EDSLs) (Axelsson 2016), and consider how to work with more than one EDSL at a time, in this case one for hardware and one for software. This paper provides methods that can be used by others who wish to combine EDSLs.

Our longer term aim is to use similar methods to program heterogeneous systems containing multicores and more than one type of accelerator (for example FPGAs and GPUs). We regard this work as a first step towards that larger goal and this is why we felt the need to develop a general approach to combining EDSLs, rather than a one-off combination of two particular EDSLs.

The paper makes the following contributions:

- We present a library for programming FPGAs in Haskell that supports hardware software co-design, generating both VHDL and C, including necessary connections between hardware and software for the Xilinx Zynq system-on-chip.
- We demonstrate the use of the library on a cryptographic example, and benchmark software and mixed hardware-software implementations. Repartitioning is far easier than it would be if the parts were written directly in C and VHDL.
- We present type-based techniques for the implementation of more than one EDSL. We build upon a generic representation of imperative programs that is loosely coupled to instruction and expression types, allowing the individual parts to be developed and improved separately.
- We implement code generation by a series of translations between progressively smaller, typed EDSLs, safeguarding against errors that arise in untyped translations.
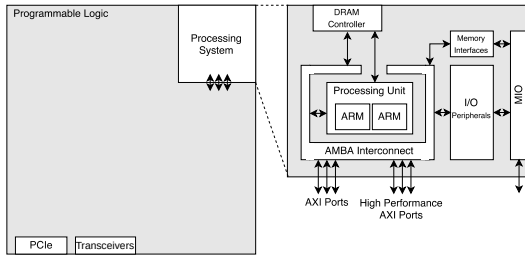
Figure 1. Overview of a Zynq system-on-chip (SoC).

2 A Co-design Language

Figure 1 shows the main parts of the Zynq system-on-chip. For our purposes here, the important parts are the ARM cores (of which we will use one, programmed in C), the programmable logic (which we will program by synthesis from VHDL) and the AXI ports that connect the two. So, starting with a single Haskell program, we have three tasks: generate C for the software parts, generate VHDL for the hardware parts, and generate a combination of C and VHDL that controls the handshaking and communication across the AXI ports to allow the software and hardware parts to communicate. We will return to this last task in section 4.

C and VHDL are both imperative. To represent them, we build on recent work on EDSLs as deep embeddings of monads (Axelsson 2016; Bracker and Gill 2014; Hickey et al. 2014; Svenningsson and Svensson 2013). The general idea is that one can view an imperative program as a sequence of instructions to be executed on some machine, which looks similar to programs written in a stateful monad. In fact, a stateful program composed with monadic operations can be directly translated into statements in an imperative language. Our hardware and software languages are both based on such a deep embedding of imperative programs:

```
data Prog ins exp a
instance Monad (Prog ins exp)
```

Having `Prog` parameterized on the instruction type, `ins`, and the expression type, `exp`, lets us define the software and hardware monads as programs, parameterized on their respective types:

```
type Software = Prog SIns SExp
type Hardware = Prog HIns HExp
```

As an example, a program for reading an integer from a stack and then putting it back again can be written in do-notation as:

```
reput :: (Monad m, Stack m) ⇒ m ()
reput = do x ← get; put x
```

`get` and `put` are brought into scope by the `Stack` constraint on the monad m, which ensures that any instantiation of m will support the two aforementioned stack operations. In the case of the stack example, software and hardware both support its operations and it can thus be interpreted in both monads.

A language embedding based on monads gives us a representation of statements in an imperative program, but most meaningful programs also include a notion of pure expressions. The expression type is, however, not readily available in the type signature of a program like reput, as it only refers to the monad m. We therefore use a type family Exp to retrieve the expression type of a program, which lets us combine instruction and expression constraints:

```
add :: (Monad m, Stack m, Num (Exp m)) ⇒ m ()
add = do x ← get; y ← get; put (x + y)
```

While `Software` and `Hardware` both support the previous stack operations, we can also point out groups of operations which are only supported by some interpretations but not by others, thus forming a hierarchy of classes. Monads form the base of the hierarchy, but operations intended for either software or hardware branches also require that their monad is an extension of their respective `Software` and `Hardware` monads. We call these monads `SoftwareM` and `HardwareM`, and they provide the language specific operations from their respective monads. For example, using the `SoftwareM` constraint, we can write a stack program that prints its top element to standard output:

```
print :: (SoftwareM m, Stack m, Printf m) ⇒ m ()
print = do x ← get; printf "%d" x
```

`print` is, according to its constraints, an entirely software based program; VHDL lacks an instruction that corresponds to the `printf` function found in C.

The interesting aspect of programming for the Zynq is that we can mix hardware and software fragments to form a whole program. As an example, we create small hardware program for the multiplication of two unsigned bytes:

```
multiply :: HExp Word8 → HExp Word8 → HExp Word8
multiply a b = a * b
```

To make the multiplier accessible from the software side, we turn it into a component called `multComp`:

```
multComp :: Component (Word8 → Word8 → Word8)
multComp = component $
    inp $ λa → inp $ λb → out $ multiply a b
```

where `inp` and `out` capture the type of the multiplier, which we normally cannot inspect in Haskell, and turn the type into a signature of the multiplier (Axelsson and Persson 2015). This signature is then encapsulated by `component` to yield a hardware program that can be invoked from software through an AXI channel or from hardware. Then, after synthesizing and placing the component on the FPGA, we get its physical address and use memory-mapped I/O to access the component from software:

```
soft :: SExp Word8 → SExp Word8 → Software ()
soft a b = do mult ← mmap "0x38C00000" multComp
              rout ← call mult (a .: b .: nil)
              printf "%d" rout
```

Here, the `mmap` function implements the necessary memory-mapping for creating a pointer to the multiplier at address "0x38C0000", from the normally virtual memory space of programs running in software on the Zynq (which commonly has Linux installed). The pointer is used by `call`, which writes its arguments a and b to the addresses of the multiplier arguments, and then listens for the result. While software interacts with a component through pointers, writes to these pointers spawn corresponding actions on the AXI channels. The AXI controller intercepts these actions and forwards them to the addressed hardware component, which results in the component, possibly, responding over AXI.

Having components and their arguments be typed does ensure some type safety when interacting with a component, assuming the correct address has been given. The manual input of addresses

comes as a consequence of having to go through the design tools and synthesize a component before we can get its physical address.

The following section presents a larger example, a software implementation of a standard function from cryptography.

2.1 Introducing the PBKDF2 example

In most applications of public-key cryptography, user security is ultimately dependent on secret texts or passwords. These secrets are often not directly applicable as keys to conventional cryptosystems, and require some processing before they can be used in cryptographic operations. Moreover, as passwords are often chosen by users, they come from a relatively small space and special care is required in that processing to defend against search attacks.

A common approach to password-based cryptography, as described by Morris and Thompson (1979), is to combine a password with a salt to produce a key. The salt can be viewed as an index into a large set of keys derived from a password, and thus needs not be kept secret. While the salt does not rule out search based attacks, it does limit an opponent to searching through passwords separately for each salt. Another common approach to password-based cryptography is to construct the keys from relatively compute intensive techniques, thereby increasing the cost of a search-based attack. One such technique for hash-based key derivations is to include an iteration count, indicating how many times the hash function should be iteratively applied to produce a key.

Salt and iteration count of a pseudorandom function form the basis of PBKDF2 (Moriarty 2017), a common key derivation technique for web-based applications. PBKDF2 applies a pseudorandom function, such as the SHA1 hash-based message authentication code (HMAC-SHA1), to an input password along with a salt. Internally, the key derivation is based around the iterative hashing of blocks of bytes that each represent a part of the derived key, where a block's length is the length in octets of the pseudorandom function's output. The derived key is then obtained from the concatenation of these blocks:

$$Key = \text{take } len \text{ of } (T_1||T_2||\ldots||T_l) \tag{1}$$

where $||$ is concatenation, l denotes the number of blocks required to represent the len bytes of the derived key, and "take len of" extracts the first len bytes of a byte string. Individual blocks of T are defined as:

$$T_i = F(P, S, c, i) \tag{2}$$

Here, P is the password, encoded as a byte string, S is the salt, also encoded as a byte string, and c the iteration count as a positive integer. The function F is defined as the exclusive-or sum of the first c iterates of the underlying pseudorandom function (HMAC-SHA1) applied to the password P and the concatenation of the salt S and the block index i:

$$F(P, S, c, i) = U_1 \oplus U_2 \oplus \ldots \oplus U_c \tag{3}$$

with individual blocks of U defined as:

$$U_1 = \text{HMAC-SHA1}(P, S||INT(i))$$
$$U_i = \text{HMAC-SHA1}(P, U_{i-1}) \tag{4}$$

where $INT(i)$ is a four-byte encoding of the integer i.

To implement PBKDF2, we use some helper functions to model its three main components: T, F, and U. Due to the combinatorial nature of these components, they can be implemented quite nicely using the standard `foldl` and `map` functions in Haskell. In the co-design library, however, we make use of the methods of resource-aware functional programming developed for the Feldspar EDSL. Memory management is done through explicit, monadic operations, giving the user control over a program's memory usage. This entails the implementation of new combinators such as:

```
foreach :: (Comput m, Typ m Int, Typ m Word8)
   ⇒ (Exp m Int → m (Arr Word8)) → Exp m Int
   → m (Arr Word8)
iterate :: (Comput m, Typ m Int, Typ m Word8)
   ⇒ (Arr Word8 → m (Arr Word8)) → Exp m Int
   → Arr Word8 → m (Arr Word8)
```

where the `Comput` constraint asserts that any interpretation of `m` will support the computation operations needed to implement the functions and `Typ m a` is a collection of constraints that ensures `a` is well typed and supports the necessary expressions:

```
type Typ m a = (Num (Exp m a), ...)
```

The first combinator, `foreach`, applies its function to each index number from zero up to the given limit and returns the concatenation of the arrays. The second combinator, `iterate`, returns an array given by iterative applications of its function to an initial array. In the setting of PBKDF2, we can use `foreach` to create each block of T and `iterate` to create and join together the blocks of U. Now, PBKDF2 can be expressed as:

```
pbkdf2 :: (Comput m, Typ m Int, Typ m Word8)
        ⇒ Arr Word8 → Arr Word8 → Exp m Int
        → Exp m Int → m (Arr Word8)
pbkdf2 pswd salt c l = take l <$> ts where
  u0 i = int4b i >>= concat salt >>= hmac pswd
  uN a = hmac pswd a >>= zipWith xor a
  f  i = u0 i >>= iterate uN (c-1)
  ts   = foreach f (1 + l `div` 20)
```

`zipWith`, `take` and `concat` are like the standard functions from Haskell, but they handle arrays instead of lists; except for `take`, which reuses memory, the result type of these functions is monadic as they use new memory when creating arrays.

The underlying HMAC-SHA1 function is given by `hmac` and `sha1`, which are implemented in a similar manner as `pbkdf2` and given the following type signatures:

```
hmac :: (Comput m, Typ m Int, Typ m Word8)
        ⇒ Arr Word8 → Arr Word8 → m (Arr Word8)
sha1 :: (Comput m, Typ m Int, Typ m Word8)
        ⇒ Arr Word8 → m (Arr Word8)
```

Interpreting and compiling `pbkdf2` as a software program for an iteration count of 4096 and 256 blocks—the wpa2 encryption standard—yields the following C code (with variable declarations and the inlined `hmac` and `sha1` functions elided to save space):

```
for (v3 = 0; v3 ≤ 256 / 20 - 1; v3++) {
    a4[0 + 0] = (uint8_t) ((uint32_t) v3 >> 24);
    a4[1 + 0] = (uint8_t) ((uint32_t) v3 >> 16);
    a4[2 + 0] = (uint8_t) ((uint32_t) v3 >> 8);
    a4[3 + 0] = (uint8_t) (uint32_t) v3;
    memcpy(a5, a1, 16 * sizeof(uint8_t));
    memcpy(a5 + 16, a4, 4 * sizeof(uint8_t));
    a6 = { hmac over a5 and password. }
```

```
for (v7 = 1; v7 ≤ 4096 - 1; v7++) {
  a8 = { hmac over a6 and password. }
  for (v9 = 0; v9 ≤ 20; v9++) {
    a6[v9 + 0] = a6[v9 + 0] ^ a8[v9 + 0];
}}
  memcpy(a2 + (v3 * 20), a6, 20 * sizeof(uint8_t));
}
```

While the software implementation of pbkdf2 can be obtained by instantiation, offloading either sha1 or hmac to hardware requires that we modify the program to make use of components. As an example, we offload sha1 to hardware. sha1 is only constrained to computational programs, and can be interpreted as a hardware program. The first step towards offloading sha1 is straightforward: wrap it in a component.

```
sha1C :: Component ([Word8] → [Word8])
sha1C = component $ inpArr 64 $ λmessage →
  outArr 20 $ sha1 message
```

Here, inpArr and outArr are similar to the signature functions inp and out shown in section 2, but they handle arrays instead, represented by lists in the signature. Note also the length argument for the input and output arrays, as any array sent over AXI must have its length known at compile time to enable construction of an equal length array on the receiving end. As we are using passwords, which are normally quite short, a length of 64 is enough.

The second step to offloading sha1 is to edit the hmac function, which currently calls the old sha1 and is defined as:

```
hmac message salt =
  do opad ← mapA (0x5c `xor`) salt
     ipad ← mapA (0x36 `xor`) salt
     fst ← sha1 =<< concatA ipad message
     sha1 =<< concatA opad fst
```

where mapA and concatA behave like their similarly named Haskell functions, but have been adapted to work with arrays instead.

From inspecting hmac, we see that sha1 is used as a regular combinator, for hashing the message twice with two different salts. As the calls to hardware components behave like regular function calls, we can simply replace each occurrence of sha1 in hmac with a function that calls the hardware component instead. For that purpose, we define callSha1,

```
callSha1 :: Arr Word8 → Software (Arr Word8)
callSha1 arr = do c ← mmap "0x38C00000" sha1C
                  call c (arr .: nil)
```

The need for an explicit hex address, as with the previous soft example, comes from the use of an external design tool to synthesize the component (and we plan to automate this process). The last two lines of hmac are updated:

```
    fst ← callSha1 =<< concatA ipad message
    callSha1 =<< concatA opad fst
```

These two changes produce a PBKDF2 design that makes use of both software and hardware.

Interpreting and compiling the new pbkdf2 as a software program, with sha1 offloaded to hardware, yields roughly the same C code as before (minus the sha1 code). The biggest change comes from the extra pointers used for sending and receiving data from the sha1 component:

```
f_map(0x83C00000, (void**) &output_sha1,
      &offset_sha1);
output_sha1 = output_sha1 + offset_sha1;
int * input0 = output_sha1 + 21;
int * output0 = output_sha1;
```

where f_map handles the memory mapping for connecting a pointer to the physical address of the component. input0 and output0 are then connected to the beginning of the component's input and output addresses, respectively. The remaining C code is then free to use them as regular variables, where writing and reading are done synchronously.

3 Implementation

In the following section, we present the techniques, types and abstractions used to implement a co-design library based on monads. While the focus is on implementing a hardware software co-design language, the ideas presented apply generally to the construction of multiple EDSLs that can be mixed. The four steps are:

1. *Implement a deeply embedded, imperative core language.* Its purpose is not to act as a convenient user interface, but rather to accurately describe the two target languages and offer an efficient platform for evaluation and code generation. As a consequence, it is kept as simple as possible.
2. *Give monad instances for the embedded languages.* A monad instance gives us access to useful combinator libraries as well as syntactic support to seamlessly mix computations from the embedded core languages with those in Haskell.
3. *Build higher-level extensions of the core languages.* The purpose of extending the core languages is to get back the user-friendly abstractions that were forsaken by the simpler core languages. Translate each extended language into its core language, to avoid costs associated with the abstraction.
4. *Implement user-friendly interfaces as a hierarchy of shallow embeddings on top of the extended core languages.* Each interface is given as a separate type-class and provides overloaded operations on that type. Structuring the type-classes in a hierarchy gives guarantees about the presence, and absence, of constructs in an embedded program.

3.1 Imperative programs

Inspired by the work of Svenningsson and Svensson (2013) and by the Operational Monad (Apfelmus 2017), we implement the first step outlined above, and declare our monadic representation of imperative programs. Like our inspiration, we capture monadic computations with an algebraic data type, but we also take into account the fact that different languages support different operations as well as expressions. We call the representation Prog:

```
data Prog ins exp a where
  Return :: a → Prog ins exp a
  Instr  :: ins (Prog ins exp) exp a
         → Prog ins exp a
  (:>>=) :: Prog ins exp a → (a → Prog ins exp b)
         → Prog ins exp b
```

In addition to having an extra parameter, our representation of programs has instructions as higher-order functors, parameterized on the program they are part of and its expression type. This facilitates an extensible definition of instructions, as shown in section 3.2.

As one would expect from a deep embedding of monads, `Prog` includes constructs to represent the standard monadic *bind* and *return* operations: `Return a` lifts a value a without introducing any effects; `m :>>= f` executes m and applies f to its result, producing a new computation as a result. The monadic instance declaration for `Prog` thus follows naturally[1]:

```
instance Monad (Prog ins exp) where
  return = Return
  (>>=)  = (:>>=)
```

The monadic instance for programs corresponds to the second step outlined above, as it provides syntactic support for sequencing instructions in either language.

The gist of the idea behind separating the instructions of `Prog` from its monadic constructs is straightforward: an instruction's effect will only depend on its interaction with other instructions, permitting their sequencing to be handled separately. The task of implementing a language based on `Prog` is then the same as writing an interpreter for the language's instructions. In particular, we can define a generic interpreter for programs that maps them to their intended meaning:

```
interpret :: Monad m
  ⇒ (∀a . ins (Prog ins exp) exp a → m a)
  → Prog ins exp a → m a
interpret f (Instr i) = f i
interpret f (Return a) = return a
interpret f (m :>>= k) =
  interpret f m >>= interpret f . k
```

`interpret` lifts a monadic interpretation of the primitive instructions, which may be of varying types, to a monadic interpretation of the whole program. By using different types for the monad m, one can implement different "back ends" for programs. For example, interpretation in Haskell's IO monad gives a way to *run programs*, while interpretation in a code generation monad can be used to make a *compiler to another language*.

The ideas behind `Prog` and `interpret` are perhaps best understood by looking at a concrete example. We define `Comp`, a collection of purely computational instructions:

```
data Comp prg exp a where
  NewRef :: Type a ⇒ exp a → Comp prg exp (Ref a)
  GetRef :: Type a ⇒ Ref a → Comp prg exp (Val a)
  SetRef :: Type a ⇒ Ref a → exp a
            → Comp prg exp ()
  For :: (Type a, Integral a) ⇒ exp a → exp a
        → (Val a → prg ()) → Comp prg exp ()
```

`NewRef`, `GetRef` and `SetRef` handle the creation and management of mutable references. `For` represents a for-loop over a specified range, with a parameterised function indicating the program to run at each iteration. The `Val` and `Ref` types play a particular rôle. `Val` is used whenever a pure value is produced by one of the instructions and represents a value in *any* expression; returning a value in exp directly would be problematic for translating between instructions with different expression types. For similar reasons, any reference returned by an instruction is represented by `Ref`.

---

[1]The observant reader may suspect that `Prog` does not obey the monad laws, since it allows us to observe the nesting of binds in a program, and indeed this is the case. In fact, we only interpret `Prog` in ways that are ignorant of the nesting of bind, but there is nothing to guarantee this.

Different representations of `Val` and `Ref` are needed in different interpretations. For instance, when an instruction involving values of type `Val a` is evaluated, a concrete value of type a is needed, while compilation needs a symbolic representation of the value. To accommodate both representations, we use a sum type:

```
data Val a = ValE a        | ValC String
data Ref a = RefE (IORef a) | RefC String
```

Using sum types for values and references is a simple solution to having these types assume different roles during interpretation[2], but those interpretations will sometimes have to do incomplete pattern matching as a result. For example, we have to assume that values are constructed with `ValE` during evaluation, as values of `ValC` cannot be evaluated. Luckily, by making `Val` and `Ref` abstract types, we can hide any such concerns from users.

The predicate type `Type` restricts the set of types that can be used in `Comp` instructions. For simplicity, we have assumed that both software and hardware support the types of `Type`, and that those types can be represented using the following class:

```
class    (Eq a, Ord a, Show a) ⇒ Type a
instance (Eq a, Ord a, Show a) ⇒ Type a
```

We are now ready to define the evaluation of programs with computational instructions as an interpretation in Haskell's IO monad:

```
run :: EvalExp exp ⇒ Prog Comp exp a → IO a
run = interpret runComp

runComp :: EvalExp exp
  ⇒ Comp (Prog Comp exp) exp a → IO a
runComp (NewRef a) = RefE <$> newIORef (evalE a)
runComp (GetRef (RefE r)) = ValE <$> readIORef r
runComp (SetRef (RefE r) a) =
  writeIORef r (evalE a)
runComp (For l u body) =
  mapM_ (run . body . ValE) [evalE l..evalE u]
```

Most of the heavy lifting is done by the `runComp` function, which defines the interpretation of each instruction. The IO references themselves, and their associated functions, are provided by the `IORef` library. Note that only `RefE` and `ValE` are used on both sides of the equation during evaluation; `RefC` and `ValC` are not supported in the standard interpretation.

Because running a program in the IO monad also involves running any expressions its instructions may contain, we have limited `run` to instructions that contain runnable expressions. This is captured by the `EvalExp` constraint that provides the function `evalE` for evaluating closed expressions. `EvalExp` is defined as follows:

```
class EvalExp exp where
  evalE :: exp a → a
```

Code generation can be defined using `interpret` as well, given a function `compToC` which translates individual instructions into C code but using a monad for code generation instead of IO:

```
compileC :: CompileCExp exp
          ⇒ Prog Comp exp a → C a
compileC = interpret compToC
```

---

[2]Having `Val` and `Ref` as sum types is sufficient for the example languages we implement in this paper; another approach could see them implemented as type families.

As with evaluation, the compiler itself is straightforward and it is the mapping from instructions to C code that does most of the work. Also, just as evaluation required its expressions to be runnable, compilation now requires that we constrain its instructions with `compileC` to expressions that can be compiled as well. Code generation through `interpret` is not limited to just C code, however. Compilation to hardware descriptions in VHDL only requires us to supply a VHDL code generation monad and a mapping from instructions to VHDL fragments:

```
compileVHDL :: CompileVHDLExp exp
            ⇒ Prog Comp exp a → VHDL a
compileVHDL = interpret compToVHDL
```

A quick summary of progress is in order. We have shown a method of embedding imperative programs using the general `Prog` type. The monadic sequencing of instructions in the embedding has a close correspondence to statements in an imperative language and, because it is parameterized on the kind of instructions it accepts, we can use it for both the software and hardware monads. So far, however, it might not look like we have come far in the pursuit of a unified hardware software co-design language—as programs have been restricted to purely computational instructions. Such computational instructions are part of both C and VHDL; we only need to add the language-specific instructions.

3.2 Software and Hardware

Introducing the `Prog` type has brought real progress. The problem has been reduced from defining a whole co-design language to defining its instruction sets and expression languages. For instructions, the computational ones we have seen so far are part of both the software and hardware languages, and should be reused in the definition of both languages. Fortunately, the problem of extending `Comp` with data types from either software or hardware is already solved in Data Types à la Carte (DTC) (Swierstra 2008). DTC introduces a type composition operator (`:+:`) that can be seen as a higher-kinded version of Haskell's `Either` type, and lets us define the instruction set of a language as a sum of smaller data types—like `Comp`.

We will use a tweaked version of (`:+:`) that has been adapted to work with our higher-order functors, such as the instruction type, and demonstrate its use by introducing three new, language-specific instruction sets:

```
data Printf prg exp a where
  Printf :: String → [Arg exp]
          → Printf prg exp a

data Process prg exp a where
  Process :: [Identifier] → prg ()
          → Process prg exp ()

data Signal prg exp a where
  NewSig :: Type a ⇒ exp a
          → Signal prg exp (Sig a)
  GetSig :: Type a ⇒ Sig a
          → Signal prg exp (exp a)
  SetSig :: Type a ⇒ Sig a → exp a
          → Signal prg exp ()
```

`Printf` adds an instruction that models the `printf` function from C and thus takes a string, which may contain formatting symbols, and a list of values to substitute said symbols with. `Process` and `Signal` introduce instructions for modeling processes and signals from VHDL. It takes a list of identifiers on which the process will trigger, and the program to run once it does trigger. `Signal` introduces instructions for managing signals, which are similar in behavior to variables, but introduce a small delay when writing to them—a signal's new value will not be immediately available when writing to them, and will retain the old value for one time step.

With the help of (`:+:`), the new instructions can be combined with those in `Comp` to form the instruction sets for the software and hardware monads.

```
type SIns = Comp :+: Printf
type HIns = Comp :+: Process :+: Signal
```

Unfortunately, introducing (`:+:`) also means that constructing instructions becomes more complicated:

```
newSig :: Type a ⇒ exp a → Prog HIns exp (Sig a)
newSig a = Instr (InjR (InjR (NewSig a)))
```

The problem is that instructions are now tagged with injections, and the ordering and number of these injections will change as the instruction set gets larger. Fortunately, DTC provides a solution to this problem as well: the (`:<:`) class provides an `inj` function that automatically handles injection of single instructions based on their instruction set type. The class defines a form of subsumption where its instances perform a linear search at the type level to find the right nesting of injections. With this class, we can define a function that automatically injects instructions into programs:

```
inject :: (i :<: ins) ⇒ i (Prog ins exp) exp a
       → Prog ins exp a
inject = Instr . inj
```

We can now define the smart constructor for signals as follows:

```
newSig :: (Type a, Signal :<: ins) ⇒ exp a
       → Prog ins exp (Sig a)
newSig = inject . NewSig
```

While instructions that only use `Sig` or `Ref` can be directly injected into programs, a bit more care has to be taken for those that use `Val` since we want to keep `Val` hidden from users, and instead use an abstract type like `exp` to hold values. We note that `exp` need only support lifting of pure values and variables in order to model the features of `Val`. A type class implements these operations:

```
class FreeExp exp where
  varE :: Type a ⇒ String → exp a
  litE :: Type a ⇒ a      → exp a
```

Using `FreeExp`, we can define a general function for converting `Val` to an abstract type of `exp` that works independently of which interpretation it is used in:

```
express :: (FreeExp exp, Type a) ⇒ Val a → exp a
express (ValC s) = varE s
express (ValE a) = litE a
```

Now we can define any constructors that employ `Val` internally, using `express` as needed:

```
getRef :: (Type a, FreeExp exp, Comp :<: ins)
       ⇒ Ref a → Prog ins exp (exp a)
getRef r = express <$> inject (GetRef r)
```

Note that smart constructors for computational instructions have become more general in their type than the previous constructors were. For example, the type of getRef says that we can read references in programs with *any* instruction set, as long as it contains Comp, that is, any instruction set of the form (... :+: Comp :+: ...). In other words, getRef is a valid function in *any language* that supports computational instructions.

3.3 Extensible Interpretation

DTC enables an extensible definition of software and hardware instructions, but compilation and evaluation are still locked to the closed domain of Comp. To extend interpretation as well, we need to abstract over the expression type and factor out the translation of instructions to a user-provided function. We do this by defining a new type class:

```
class Interp ins prg exp m where
  interp :: ins prg exp a → m a
```

interp describes a translation from an instruction of type ins, parameterized by its program type prg and expression types exp, into a monadic expression of type m. And, by providing an instance of Interp for (:+:), we can further separate the translation into that of single of instructions. We define the new interpreter as:

```
interpret₂
  :: (Interp ins (Prog ins exp) exp m, Monad m)
  ⇒ Prog ins exp a → m a
interpret₂ = interpret interp
```

In order to get back support for evaluation or compilation of Comp under the new interpretation scheme, we simply add an instance for it with Interp and the desired monad. The instances themselves are similar to their earlier, corresponding interpretation functions, and we can even reuse their definitions for the new instances—assuming that we rewrite the recursive interpretation of programs in for-loops to use interpret₂ instead of run:

```
instance ( Interp ins  (Prog ins exp) exp IO
         , EvalExp exp)
         ⇒ Interp Comp (Prog ins exp) exp IO
  where interp (For l u body) = mapM_
    (interpret₂ . body . ValE)
    [evalE l..evalE u]
```

The other cases of interp are defined in the same way as they were in the earlier runComp function.

3.4 Type classes

Some instructions are only meaningful in certain language interpretations. The co-design library is therefore structured with type classes. For example, a hardware program can deal with signals and processes, whereas a software program has no notion of such concepts at all. We can point out groups of instructions that are supported by some interpretations but not others, forming a hierarchy of classes. The base class is the standard Monad class from Haskell, ensuring that any language can handle the necessary information

plumbing needed for sequencing of instructions in a program. Sub-classes of Monad include References, a class that provides support for mutable references:

```
class Monad m ⇒ References m where
  newRef :: Type a ⇒ Exp m a → m (Ref a)
  getRef :: Type a ⇒ Ref a → m (Exp m a)
  setRef :: Type a ⇒ Ref a → Exp m a → m ()
```

Reference instructions, like most instructions, refer to the expression language associated with the monad m. We therefore use a type family Exp to retrieve the expression type of a program. Exp is defined as follows:

```
type family Exp m where Exp (Prog ins exp) = exp
```

Also, we should note that classes like References bear a resemblance to how programs are represented in a finally tag-less setting; a class interface gives part of the language's syntax, although an instruction's semantics are given by the appropriate Interp instances rather than instances of a class like References.

With the help of the earlier inject function for lifting expressions into program stubs, we can provide a default References instance for programs whose instructions contain Comp.

```
instance (Comp :<: ins, FreeExp exp)
    ⇒ References (Prog ins exp) where
  newRef v   = inject (NewRef v)
  getRef r   = express <$> inject (GetRef r)
  setRef r v = inject (SetRef r v)
```

Constraints like References form the fourth step outlined in section 3, and give a guarantee about the presence, or even absence, of certain instructions in a program. This, in turn, means that we can determine from a program's type whether it can be implemented in software or hardware, or both. A group of common instructions, such as the references and basic control structures found in both software and hardware, can be defined as a collection of constraints:

```
type Comput m =
  (Monad m, References m, Arrays m, Control m)
```

Programs that can only be interpreted in software will have the SoftwareM m constraint, while programs that can only be interpreted as hardware will have the HardwareM m constraint.

```
class Monad m ⇒ SoftwareM m where
  liftS :: Prog SIns (Exp m) a → m a

class Monad m ⇒ HardwareM m where
  liftH :: Prog HIns (Exp m) a → m a
```

Sub-classes of SoftwareM and HardwareM include classes that model the earlier printf instruction, and the process and signal instructions, respectively.

```
class SoftwareM m ⇒ Printf m where
  printf :: String → [Arg (Exp m)] → m ()

class HardwareM m ⇒ Process m where
  process :: [Identifier] → m () → m ()

class HardwareM m ⇒ Signal m where
  newSig :: Type a ⇒ Exp m a → m (Sig a)
  getSig :: Type a ⇒ Sig a → m (Exp m a)
  setSig :: Type a ⇒ Sig a → Exp m a → m ()
```
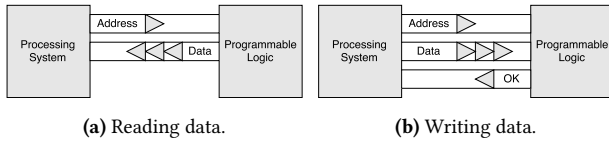
(a) Reading data. **(b)** Writing data.

Figure 2. Transactions.

A program will typically be constrained in the type to indicate what interpretations are allowed to run the program. The following program can be run by any language that supports computational operations:

```
modifyRef :: (Comput m, Type a)
  ⇒ (Exp m a → Exp m a) → Ref a → m ()
modifyRef f r = getRef r >>= setRef r . f
```

The architecture of the class system makes it easy for the user to add new classes of operations to the hierarchy, and, in combination with DTC, new interpretations that give semantics to them. As such, we have addressed the fourth point in the outline of our method at the beginning of this section. Software and hardware languages have, however, been kept separate from each other. In most co-design settings, there is interleaved execution of software and hardware programs, which therefore need to communicate with each other. The following section introduces the necessary communication channels.

4 Co-Design

The Zynq SoC couples a processing system of two ARM cores with a programmable logic array. The processing system and the programmable logic are connected by AXI4 compliant interfaces.

Full AXI4 offers a range of interconnects that include variable data and address bus widths, high bandwidth burst and cached transfers, and various other transaction features. The full specification provides great flexibility to users, and we plan to implement it eventually. Here, we make use of a simpler interconnect, AXI4-lite, which is also provided by the Zynq. This subset of the full specification forgoes the more advanced burst and streaming transactions for a simpler communication model of writing and reading single pieces of data, one at a time. Five channels make up the whole AXI4-lite specification: the read and write address channels, the read and write data channels, and the write acknowledge channel. We introduce these channels from the processor's perspective.

The read address channel allows the processor to indicate that it wishes to start reading data from a hardware component, and carries the address and handshaking information necessary to initiate the transaction. The hardware component then answers through the read data channel, which carries the data values that are sent during the transaction, along with the necessary handshaking. It is these two channels that together implement the protocol for a read data transaction. In the opposite direction, from processor to hardware, it is the write address channel that initiates a write transaction, with data flowing on the write data channel. A write response channel lets the component acknowledge a successful write or an error. Read and write transactions are visualized in Figure 2.

AXI4-lite channels are represented in a hardware description language as signals, driven by processes that implement the associated handshaking. Other than processes and signals, we also need to wrap these AXI4-lite channels in a hardware component, like the

one that we saw very briefly in section 2. Components are ordinary hardware programs, but encapsulated and given a signature which describes how software programs or other hardware blocks should connect their pointers or wires to ports on the component. We define a new hardware instruction for components:

```
data Architecture prg exp a where
  NewComp :: Signature a
    → Architecture prg exp (Component a)
  PortMap :: HArgs a → Component a
    → Architecture prg exp ()
```

The NewComp constructor deals with the creation of hardware components from programs, while PortMap is used to instantiate components, using the portmap function in VHDL to connect components. The behavior of portmap is similar to function application on the software side; we name the component to be instantiated and its arguments, but we also give the target for its output. While function application in Haskell is simply another sequential operation, portmap is a concurrent statement and is more akin to plugging a hardware component into a socket on a board.

Normally, one cannot inspect the types of a component, and thus neither its incoming nor outgoing signals, from within a program in Haskell. To ensure that any application of the component on the software side is type correct, a component's type signature is encapsulated by the Signature type (Axelsson and Persson 2015). Such signatures enable the creation of a safe mapping of each argument to its signal in the hardware component, and one can think of them as adding a top-level lambda abstraction to the hardware language. On the software side, we also make use of the Args type, as a heterogenous list of arguments to match the signature of a hardware component when calling.

We add components to the set of hardware instructions:

```
type HIns = ... :+: Architecture
```

and note that the earlier compiler no longer supports the resulting hardware programs, as it cannot interpret components yet. But, as both constructs from the component instructions are based on existing VHDL constructs, their interpretation is a straightforward translation.

Now, to implement the AXI4-lite protocol we simply need to realize the *ready* and *valid* principle that it's based on; *ready* indicates that a hardware component is ready to accept a transfer of data or an address, and *valid* signals that any data or addresses sent by the processor are valid and can be sampled. In fact, reading and writing over AXI both follow this pattern, and can both be implemented using a process with a few conditional statements and logical expressions. Other than handshaking, we also need processes for managing the routing of data to and from AXI4-lite wires to a hardware component. These extra processes can be implemented in very much the same way as the previous ones: they are clocked processes that listen on their relevant signals and, given that the correct handshaking signals have been set, load the given address or store the received data in the addressed register. It is these registers that are connected to a hardware component using the portmap function, and this allows data received through the AXI4-lite signals to be forwarded to the component.

While implementing the AXI4-lite protocol was quite a low-level ordeal, in that it mostly consists of reading and writing to bit flags, the result is a hardware function that automatically connects a given

hardware component to an AXI4-lite interconnect by inspecting its signature (we omitted parts of its signature to save space as the AXI4-lite protocol consists of over twenty signals):

```
axi4lite :: HardwareM m ⇒ Component a
    → m (Component (Signal Word32 → ...))
```

Having connected a hardware component to an AXI4-lite interconnect, the next step to get it onto the programmable logic is to compile it. The Haskell description of the component generates VHDL, which is input to a synthesis tool that produces the bit stream that causes the programmable logic array to be configured to implement the required behavior. In addition, the synthesizer gives the processor access to the inputs and outputs of the component through memory mapped I/O. As a consequence of running Linux on the processors, such memory mapping is handled by the mmap function, which creates a new mapping in the virtual address space of the calling processor. We introduce a software instruction to model memory mapping:

```
data MMap exp a where
    MMap :: String → Component a
         → MMap prg exp (Addr a)
    Call :: Addr a → SArgs (Argument a)
         → MMap prg exp (Result a)
```

MMap takes a hardware component along with its physical address on the programmable logic, gives the component an AXI4-lite wrapper using the earlier axi4lite function, and produces a memory reference in software that can be accessed using Call. Call is similar to portmap from hardware, but constructs its argument list using software expressions, and its result is the result of the signature.

We extend the previous set of software instructions to include memory mapped I/O:

```
type SIns = ... :+: MMap
```

A program that supports memory mapped I/O can freely communicate with a hardware component placed on the FPGA.

5 Expressions

Programs and their instructions represent the statements in an imperative language, but an expression language is also needed:

```
data CExp a where
    Var :: Type a ⇒ String → CExp a
    Lit :: Type a ⇒ a → CExp a
    Add :: (Num a, Type a)
        ⇒ CExp a → CExp a → CExp a
```

CExp is a type of numerical expressions with variables, integer literals and addition. By using a generalized algebraic datatype we make sure that only well-typed expressions can be constructed—barring the fact that variables can have any type.

The expression type of CExp provides little to no abstractions; its operations all have a relatively small semantic step to their corresponding operations in the target domains. Interpretation of CExp is consequently a straightforward translation of its constructs into matching expressions of the target domains. For example,

```
instance EvalExp CExp where
    evalE (Lit a)   = a
    evalE (Add a b) = evalE a + evalE b
```

It is due to these attributes that CExp forms the core software language in our library (which constitutes the first step in section 3):

```
type SoftwareC = Prog SIns CExp
```

One of the hallmarks of functional programming is the presence of powerful abstractions that hide many of the details of mundane operations, such as the sharing of values through let-bindings. Even though CExp can be translated into efficient code, from a user's perspective, it is really too simple to be useful. One way to introduce some user-friendly abstractions is to make an extension of CExp where such constructs are present. We implement the third step outlined in section 3, extending the software language, and define EExp:

```
data EExp a where ...
    Let :: Type a
        ⇒ EExp a → (EExp a → EExp b) → EExp b
```

In addition to the numerical constructs already found in CExp, EExp introduces let-binding as a new construct—using higher-order abstract syntax as a convenient technique for introducing variable binders.

```
type Software = Prog SIns EExp
```

Now, we can write a pure function that shares the result of another pure function f across an addition

```
*> let ex = λ x → Let (f x) (λy → y `Add` y)
*> :t ex
ex :: (Type a, Num a) ⇒ EExp a → EExp a
```

which resembles an expression we might have written in ordinary Haskell, barring the use of constructors to build the expression rather than the overloaded operators from Num. This discrepancy can however be remedied by declaring a Num instance for EExp. For Haskell classes or functions that cannot be instantiated by expressions, we instead provide a custom class that is intended to override its Haskell version. For example, let-bindings are built-in Haskell functions, and cannot be easily overridden. We therefore provide a type class which models let-bindings:

```
class Let exp where
    share :: (Type a, Type b) ⇒ exp a
          → (exp a → exp b) → exp b
```

In contrast to any construct of CExp, the let-binding found in EExp is a higher-level abstraction that does not have a corresponding construct in the expressions of either the software or hardware languages. As a consequence, EExp is semantically further away from its target domains and its interpretation is no longer a straightforward translation. However, rather than providing a new interpretation for EExp, which is tedious and error-prone task, we define a translation to programs of the simpler CExp type, which already supports evaluation and compilation. The interesting case occurs during the translation of Let. Since let-bindings cannot be represented in CExp, we instead have to realize them through computational constructs. We provide one possible translation, translating a let-binding into reference statements:

```
elaborateSoft :: EExp a → SoftwareC (CExp a)
elaborateSoft (Let v body) = do
  r ← newRef =<< elaborateSoft v
  x ← inject (GetRef r)
  elaborateSoft $ body $ express x
```

The argument expression of `Let` is evaluated first, and the resulting value is then bound to a variable and fed into the function. Note the explicit use of the `GetRef` construct instead of its constructor function. This is because the constructor is instantiated for the core language and therefore expects a `CExp` expression, but the function we wish to pass the variable into expects an `EExp` expression.

At this point, we have managed to define both the core languages and the extended languages for software and hardware. For the core languages, we have a way of evaluating and generating source code from programs. On the other hand, for the extended languages we have the ability to write high-level programs. What is missing before we can evaluate or compile the high-level programs is a function that takes either language's elaboration function and lifts it to an interpretation of entire programs. That is, and we generalize the problem a bit here, we need a way to elaborate the expressions of one program type to a program over another expression type. Fortunately, programs are regular monads like any other, and we can actually implement this interpretation using $interpret_2$:

```
elaborateS :: Software a → SoftwareC a
elaborateS = interpret₂
```

We can now define a compiler and an evaluator for software programs by piecing together the elaboration of expressions and the corresponding interpreter for core programs.

```
runS :: Software a → IO a
runS = runSoft . elaborateS

compileS :: Software a → String
compileS = compileSoft . elaborateS
```

To extend the running and compilation technique to hardware programs, we only need to define a similar elaboration process, but for hardware expressions instead.

The above approach to expressions has several advantages: the translation is typed, which rules out many potential errors, and is easier to write than a complete translation into source code; the low-level expression type is reusable, and can be shared as an elaboration target between multiple high-level expression types. Each stage opens up for different optimization possibilities—higher-level expressions reveal semantic information that's not readily available in the lower-level languages. Separating the two language versions also, in turn, separates the compiler from the user interface, allowing the two to be improved independently of each other.

6 Evaluation

We aim for high performance code in both hardware and software (that is VHDL and C). We have not yet done extensive benchmarking, but the generated code typically performs well in both C and VHDL, being comparable to handwritten implementations. Compared to optimized code, like the C implementation of SHA1 from OpenSSL (The OpenSSL Project 2003), our naive implementation is slower by a factor of two for smaller messages. As the library programmer has access to C and VHDL at a relatively low level,
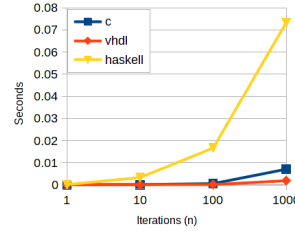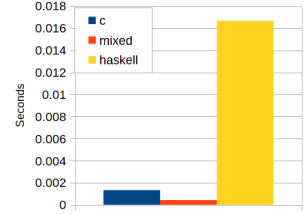

Figure 3. SHA1

Figure 4. PBKDF2

it would, however, be feasible for our users to implement similar optimizations. Still, both our C and VHDL implementations show a definite advantage over their corresponding Haskell implementations, as show in Figure 3. The three implementations of SHA1 were run on a Zynq SOC (FPGA with an ARM A9 core), where the C and Haskell versions ran on the ARM core and the VHDL component on the FPGA, but reading and writing values over an AXI channel to a software test bench. The test bench consisted of iteratively applied the three different hashes a number of times to an initial password. The hardware and software programs perform equally well, up until a thousand iterations where the hardware program performs slightly better. The reference Haskell implementation performs quite poorly in comparison.

We have also evaluated how well the library handles the bigger PBKDF2 example from section 2 in two variants, one that runs entirely in software (C) and another that offloads SHA1 to the FPGA and runs its remaining HMAC and PBKDF2 in software (Mixed). The performance of the two designs is compared in figure 4, along with a comparison to the reference implementation in Haskell (Compiled with GHC 8.0.2 with the `-O2` flag). These initial results show promising performance.

7 Related Work

There has been much work on FPGA design in functional programming languages. Some approaches concentrate on describing hardware *structure* (Baaij et al. 2010; Bachrach et al. 2012; Bjesse et al. 1998; Claessen et al. 2003; Gill et al. 2010), while others describe *semantics* and synthesise hardware from more general functional programs (Ghica et al. 2011; Zhai et al. 2015), see the survey by Gammie (Gammie 2013). SPIRAL (Puschel et al. 2005) demonstrates the advantages of using a Domain Specific Language equipped with algebraic laws in the generation of either hardware or software. However, none of these approaches, to our knowledge, considers the problem of how to develop the software and the hardware of a system implemented on FPGA *together*, although we note that a Master's thesis describes a first attempt at a foreign function interface for CλaSH to enable offloading of tasks from software to hardware (Vossen 2016).

Bluespec's BSV is translated to synthesisable Register Transfer Level (RTL) code (Nikhil 2004). With its facilities for expressing concurrency and its libraries for generating finite state machines, pipelined architectures and much else, it provides an attractive alternative for FPGA design using a higher level language. Still, as far as we know, it concentrates on the hardware side, linking to standard synthesis tools.

George et al. (2013) explore the use of Lightweight Modular Staging (LMS) to ease the construction of a domain-specific High Level Synthesis (HLS) system. They argue that this eases the reuse of modules between different HLS flows, and makes it easier to link to existing tools such as C-to-RTL compilers and core generators. Though the language-specific challenges for LMS are different from ours, the two approaches are comparable in terms of capability. The way in which code generation of our programs is built upon the idea of monadic reification is also reminiscent of Sunroof (Bracker and Gill 2014), a DSL for generating JavaScript.

Outside the domain of HLS, examples of DSLs with similar ambitions to ours include the Cryptol DSL for cryptography (Browning and Weaver 2010) and Microsoft's Accelerator (Tarditi et al. 2005) for programming GPUs and various other platforms. In both DSLs, loops are automatically unrolled during code generation. Our library thus provides a step up in expressiveness, as its keeps loops in its generated code. On the other hand, Cryptol and Accelerator provide language constructs that are well suited for expressing data-flow algorithms, at a level of abstraction appropriate for hardware—although the latest versions of Cryptol no longer support hardware generation.

Early work by Mycroft and Sharp, based on the resource aware functional programming language SAFL, explored partitioning of SAFL code into software and hardware parts, with the software parts running on generated soft processors also on the FPGA (Mycroft and Sharp 2001). IBM's Liquid Metal aimed to use a single programming language, lime, based on Java, to program FPGAs, GPUs and multicores (Auerbach et al. 2012). However, this ambitious project seems to have been cancelled, and the final publicly released version supports only FPGAs and CPUS (and not GPUs). It does not support driving vendor synthesis tools or execution on a real FPGA board.

The closest work to ours that we have found is Dave's work on Bluespec Codesign Language (BCL), an extension of Bluespec that is a unified language model (based on guarded atomic actions) for hardware-software codesign (Dave 2011). A notion of computational domain resembling an extra type abstraction is introduced to express partitioning into hardware and software. Synchronizers, which are often variants of a FIFO interface, implement the communication between hardware and software parts. The thesis concentrates on formal reasoning about partitions and refinements, and so there is no description or evaluation of examples actually running on FPGA. A later MIT thesis made further progress towards implementing BCL by embedding it in Haskell (King 2013) and an impressive study of various partitionings of a set of benchmarks is performed.

Our approach lacks the notion of rewrite rule that BCL borrows from Bluespec and it will be interesting to see if we can do without this expressive feature. One option that we should consider is using Bluespec rather than VHDL as the hardware description language. On the software side, we benefit from building upon a recent set of tools for EDSL implementation in Haskell (Axelsson 2016).

Our current implementation demands that the user program at a relatively low level of abstraction (for both software and hardware). In order to make better use of the embedding in Haskell, we would need to add layers that provide Lava-like combinators to express common regular circuit patterns (Bjesse et al. 1998), behavioural constructs like those implemented in York Lava (Naylor et al. 2009) and a synchronous programming layer that makes use of our earlier work on streams in embedded languages (Aronsson et al. 2015). Hopefully, raising the level of abstraction will allow us to protect the user from having to think about the details (such as explicit memory mapped addresses) of the communication between hardware and software.

We would also like to take advantage of having a single description of the entire system by providing the means to *test* the resulting hardware and software, for example by building on work on BlueCheck (Naylor and Moore 2015), which is a version of QuickCheck for BSV. We feel that it would be an advantage to have a single language for expressing properties of both software and hardware, but we need to perform experiments to verify this. We would also like to provide facilities for formally verifying the hardware using SAT-based techniques, as in Chalmers Lava. These are our intended next steps.

Teich (2012) provides an excellent survey on hardware software codesign in general, as well as some predictions for the future. In particular, according to Teich: "too little effort is spent in this important area of joint coverification of hardware and software". We would like to apply property-based testing here.

8 Discussion

We have presented a hardware software co-design language embedded in Haskell. The library introduces Prog as an extensible representation of imperative programs, and we showed the basic methods for designing a hardware and a software language with it. Together, the two imperative languages form the basis of the co-design library. The library allows a single program description to be interpreted in many different ways, so that evaluation in Haskell's IO monad or compilation to hardware and software are supported. Furthermore, new instructions and interpretations can be added with relatively little disturbance to the existing system, allowing the library to be used in the development of new languages. Even if software and hardware end up as the only embedded languages, we feel that the idea of establishing an interface towards the interpreter and classes for compilation and evaluation is a good one; keeping them separated lets us treat the languages and their interpreters separately. To be able to provide these features, we have relied on Haskell's type system, including monads and type classes, and higher-order functions for capturing patterns. For example, the interpret function shown earlier is a higher-order function that takes a function for interpreting single instructions, and gives us an interpretation for entire programs. So, with the help of a type class, users wishing to define a new interpretation only have to give an instance that describes how to interpret their instructions. They get the rest for free.

We have shown how to generate C and VHDL from the imperative software and hardware programs using an extensible compiler. We could have chosen to target a high level language, such as OpenCL or SystemC, to describe both software and hardware. In many applications, this is an attractive approach, as it permits software developers to do hardware design without extensive training. However, for the kinds of high performance acceleration tasks that we aim to support, this does not seem to work well (see recent experiments with OpenCL on Zynq by Svensson (Svensson 2017)). Also, recent work by (Zohouri et al. 2016) showed that it was difficult to get good performance on FPGA, compared to GPUs for instance, for HPC kernels using OpenCL, even with specialized, FPGA-oriented optimization of the OpenCL code.

We have tried to keep the translation between design and source code transparent, for example with explicit memory management, avoiding high level synthesis or compilation steps that are difficult to control or predict. We believe in leaving matters firmly in the hands of the programmer. It would be significantly harder to analyze the performance or memory profile of programs were we to rely on high level synthesis. Later, we expect to introduce autotuning, but again giving control of the search for good solutions to the programmer, in the style of Vollmer et al. (2015).

To remain focused on reusable techniques, we have left out or simplified important details of the full implementation of the co-design language. The real system[3] makes use of a sophisticated software stack supporting the development of EDSLs in Haskell, with a rich set of expressions and types. In particular, we use Axelsson's `imperative-edsl`[4] library for C generation, `hardware-edsl`[5] for VHDL generation, and a standalone package for imperative programs called `operational-edsl`[6], also developed with Axelsson.

The final system will support the full AXI protocol, to permit streaming computations and to gain access to the high performance AXI ports. We are, therefore, in the process of implementing it, making use of the library's own VHDL implementation. We plan to introduce streaming not as another primitive instruction in the language, but rather as a layer on top of the current library, where it will extend the current software and hardware programs with support for expressing synchronous data-flow. Another possible extension would be to provide access to the remaining components of the Zynq, such as its DSP slices or other processing cores, without having to rely on the synthesis tool. Such extensions would require us to introduce new layers with syntactical support for more a fine-grained execution control than simply hardware or software. Its would also be interesting to explore the opposite direction and introduce an actor language for describing control flow at higher-level in the style of Bluespec (Nikhil 2004).

Writing the co-design library has been an educational exercise in the engineering of software and hardware programs. To make the library more usable, we need to explore additional interpretations. For example, we would like to explore parallelism for software programs and to introduce verification of programs as an additional interpretation. Nevertheless, the library is an interesting practical application of Haskell to the area of hardware software co-design. As was evident from the implementation of the PBKDF2 cryptographic function, the library allows different hardware software partitions to be easily explored. It is this ease of design exploration that will likely prove to be the key advantage of this approach to co-design, especially when we explore greater uses of meta-programming. Our experience from Lava indicates that this will be a fruitful area for further research.

Acknowledgements

References

H. Apfelmus. 2017. The Operational Monad Tutorial (Blog Post). http://apfelmus.nfshost.com/articles/operational-monad.html. (2017).

---

[3] https://github.com/markus-git/co-feldspar
[4] https://github.com/emilaxelsson/imperative-edsl
[5] https://github.com/markus-git/hardware-edsl
[6] https://github.com/emilaxelsson/operational-edsl

Markus Aronsson, Emil Axelsson, and Mary Sheeran. 2015. Stream Processing for Embedded Domain Specific Languages. In *Revised Selected Papers from 26th Int. Symp. on Implementation and Application of Functional Languages (IFL)*. ACM.

Joshua Auerbach et al. 2012. A Compiler and Runtime for Heterogeneous Computing. In *Proc. 49th Design Automation Conference (DAC)*. ACM.

E. Axelsson. 2016. Compilation as a Typed EDSL-to-EDSL Transformation (Blog post). http://fun-discoveries.blogspot.se/2016/03/. (2016).

E. Axelsson and A. Persson. 2015. Programmable Signatures. In *Trends in Functional Programming, Revised Selected Papers*. Springer.

Christiaan Baaij, Matthijs Kooijman, Jan Kuper, Arjan Boeijink, and Marco Gerards. 2010. CλaSH: structural descriptions of synchronous hardware using Haskell. In *Proc. 13th Euromicro Conference on Digital System Design: Architectures, Methods and Tools (DSD)*. IEEE.

Jonathan Bachrach et al. 2012. Chisel: Constructing Hardware in a Scala Embedded Language. In *Proc. 49th Design Automation Conference (DAC)*. ACM.

P. Bjesse, K. Claessen, M. Sheeran, and S. Singh. 1998. Lava: Hardware Design in Haskell. In *Proc. Third Int. Conf. on Functional Programming (ICFP)*. ACM.

Jan Bracker and Andy Gill. 2014. Sunroof: A Monadic DSL for Generating JavaScript. In *Proc. 16th Int. Symp. on Practical Aspects of Declarative Languages*. Springer International Publishing, 65–80.

Sally Browning and Philip Weaver. 2010. Designing tunable, verifiable cryptographic hardware using Cryptol. In *Design and Verification of Microprocessor Systems for High-Assurance Applications*. Springer, 89–143.

K. Claessen, M. Sheeran, and S. Singh. 2003. Using Lava to design and verify recursive and periodic sorters. *IJSTTT* 4, 3 (2003).

Nirav Dave. 2011. *A Unified Model for Hardware/Software Codesign*. Ph.D. Dissertation. EECS Dept., MIT.

Peter Gammie. 2013. Synchronous Digital Circuits As Functional Programs. *ACM Comput. Surv.* 46, 2, Article 21 (Nov. 2013), 27 pages.

N. George, D. Novo, T. Rompf, M. Odersky, and P. Ienne. 2013. Making domain-specific hardware synthesis tools cost-efficient. In *2013 International Conference on Field-Programmable Technology (FPT)*. 120–127.

Dan R. Ghica, Alex Smith, and Satnam Singh. 2011. Geometry of Synthesis IV: Compiling Affine Recursion into Static Hardware. In *Proc. 16th Int. Conf. on Functional Programming (ICFP)*. ACM.

Andy Gill et al. 2010. Introducing Kansas Lava. In *Proc. 21st Int. Conf. on Implementation and Application of Functional Languages (IFL)*. Springer-Verlag.

Patrick C. Hickey, Lee Pike, Trevor Elliott, James Bielman, and John Launchbury. 2014. Building Embedded Systems with Embedded DSLs (Experience Report). In *Proc. 19th Int. Conf. on Functional Programming (ICFP)*. ACM.

Myron King. 2013. *A Methodology for Hardware-Software Codesign*. Ph.D. Dissertation. EECS Dept., MIT.

K. Moriarty. 2017. Password-Based Cryptography Specification Version 2.1. (2017). https://tools.ietf.org/html/rfc2898

Robert Morris and Ken Thompson. 1979. Password Security: A Case History. *Commun. ACM* 22, 11 (Nov. 1979), 594–597.

Alan Mycroft and Richard Sharp. 2001. *Hardware/Software Co-design Using Functional Languages*. Springer-Verlag, 236–251.

M. Naylor and S. Moore. 2015. A generic synthesisable test bench. In *ACM/IEEE Int. Conf. on Formal Methods and Models for Codesign (MEMOCODE)*. 128–137.

Matthew Naylor, Colin Runciman, and Jason Reich. 2009. The Reduceron home page, fetched May 2017. (2009). https://www.cs.york.ac.uk/fp/reduceron/

Rishiyur Nikhil. 2004. Bluespec System Verilog: efficient, correct RTL from high level specifications. In *ACM/IEEE Int. Conf. on Formal Methods and Models for Co-Design (MEMOCODE)*. IEEE, 69–70.

M. Puschel et al. 2005. SPIRAL: Code Generation for DSP Transforms. In *Proc. IEEE*, Vol. 93. Issue 2.

Josef David Svenningsson and Bo Joel Svensson. 2013. Simple and Compositional Reification of Monadic Embedded Languages. *Proc. 18th Int. Conf. on Functional Programming (ICFP)* (2013).

Bo Joel Svensson. 2017. OpenCL Reduction on the ZYNQ. (2017). http://svenssonjoel.github.io/writing/zynqreduce.pdf

Wouter Swierstra. 2008. Data Types à La Carte. *J. Funct. Program.* 18, 4 (July 2008).

David Tarditi, Sidd Puri, and Jose Oglesby. 2005. Accelerator: simplified programming of graphics processing units for general-purpose uses via data-parallelism. *Rapport Technique, Microsoft Research* (2005).

J Teich. 2012. Hardware Software Codesign: the Past, the Present, and Predicting the Future. *Proc. of the IEEE* 100 (2012).

The OpenSSL Project. 2003. OpenSSL: The Open Source toolkit for SSL/TLS. (April 2003). www.openssl.org.

Michael Vollmer, Bo Joel Svensson, Eric Holk, and Ryan R. Newton. 2015. Meta-programming and Auto-tuning in the Search for High Performance GPU Code. In *Proc. 4th Int. Workshop on Functional High-Performance Computing (FHPC)*. ACM.

J.J. Vossen. 2016. *Offloading Haskell functions onto an FPGA*. Master's thesis. Univ. Twente.

Kuangya Zhai, Richard Townsend, Lianne Lairmore, Martha A. Kim, and Stephen A. Edwards. 2015. Hardware Synthesis from a Recursive Functional Language. In *Proc. 10th Int. Conf. on Hardware/Software Codesign and System Synthesis (CODES)*. IEEE.

Hamid Reza Zohouri, Naoya Maruyama, Satoshi Matsuoka, and Aaron Smith. 2016. Evaluating and Optimizing OpenCL Kernels for High Performance Computing with FPGAs. In *Proc. SuperComputing Conference*. IEEE.